

EIKI

...

# ① What is AI?

→ No easy answer to what is or isn't AI

Reasons: → no official definition

→ Science fiction influence

→ Easy tasks can become very hard

→ Hard tasks are mastered on a non-human level

→ Public perception of AI is nebulous ("nebig", "schleierhaft")

The science and engineering of **making intelligent machines**, especially intelligent computer programs. It is related to the similar task of using computers to **understand human intelligence**, but AI does not have to confine itself to methods that are biologically observable.

John McCarthy, Turing Award 1971

The science of making machines do things that would require **intelligence** if done by men.

Marvin Minsky, Turing Award 1969

## What is Intelligence

Turing Test: When does a system behave intelligently?

Assumption → if one entity cannot be distinguished from another by observing its behavior.

Test → 1. Human interrogator interacts with two players, A and B (one being a computer) blind.

2. If interrogator cannot determine whether A or B is a computer or human, the computer passes the test.

→ Test still relevant today → different question → requires collaboration of major components of AI: knowledge, reasoning, language understanding, learning, ...

But hard/not reproducible and constructive or amenable to mathematical analysis.

Followup Question: Does being human-like mean you are intelligent?

Chinese Room Argument: Is intelligence the same as intelligent behavior?

Assumption → Even if a machine behaves in an intelligent manner, it does not have to be intelligent at all

Test → 1. Person who doesn't know Chinese locked in a room. Outside person with paper, who can only communicate by writing Chinese and sliding notes under door.

2. Person inside has detailed instructions how to answer in Chinese but can't translate it or understand it.

Follow-up Question: Is person in room intelligent? Is a self-driving car intelligent?

The question of how an intelligent being is defined, is a long and difficult one, but...

Characteristics of AI → One AI to rule them all?

"unlikely to have any more effect on the practice of AI research than philosophy of science generally has on the practice of science."

John McCarthy

General vs Narrow AI: General / Strong AI → can handle any intellectual task

↳ Same as being vs. acting. Narrow/weak AI → specified to deal with a concrete intelligent OR a set of specified tasks.

→ Strong AI is goal in research, currently we are primarily using narrow AI.

Adaptability: improve performance by learning from experience

Autonomous: perform tasks in environments without constant guidance by user/expert

Rationality: Law of Thoughts → beginning of reasoning and the question of what 'correct' argument and thought processes are.

Rational Behavior → question of "doing the right thing". In systems that act rationally → "right thing" is what is expected to maximize the achievement given available information

→ Rational behavior has two advantages over law of thoughts → more general (in many situations, provably correct action non-existent)  
more amenable (rationality can be defined and optimized)

However rarely a very good model of reality

Different definitions of AI due to different criteria: Two dimensions → 1. Thought process/reasonings vs. behavior /action

2. Success according to human standards vs. to an idea/concept of intelligence(rationality)

	Human	Rationality
Think	Systems that think like humans	Systems that think rationally
Act	Systems that act like humans	Systems that act rationally

Cognitive Science: Theory of human thinking based on scientific theories of internal brain activities (cognitivist mode): → predicting and testing human behavior  
(systems that think like humans) → identification from neurological data

↳ Brings together computational models from AI and experimental techniques from psychology to construct precise and testable theories of the mind.  
Often viewed distinct from AI → wrong!

Cognitive Neuroscience: How does brain work at neuronal level?

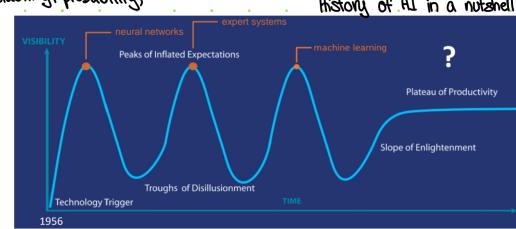
**Foundations of AI:** Different fields have contributed to AI and build foundation of modern AI →

- Philosophy (logic, reasoning, mind as physical system, foundations of learning, language and rationality)
- Mathematics (formal representation and proof algorithms, computation, (un-)decidability, (in-)tractability, probability)
- Psychology (adaptation, phenomena of perception and motor control)
- Economics (formal theory of rational decisions, game theory)
- Linguistics (knowledge representation, grammar)
- Neuroscience (physical substrate for mental activities)
- Control theory (homeostatic systems, stability, optimal agent design)

Today: Nobel Prize in Physics (Hinton & Hinton → "for foundational discoveries and inventions that enable machine learning with artificial neural networks")

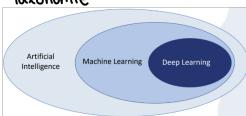
Nobel Prize in Chemistry (Baker → "for computational protein design",

Hassabis & Jumper → "for protein structure prediction")



**Related Fields and Taxonomy of AI:**

**Taxonomy**



Subdisciplines → Machine Learning, Deep learning, Search and optimization, Robotics, Natural Language Processing, Computer Vision, Cognitive Science, AI Systems, Data Science ...

**Applied AI in modern fields:** Autonomous Planning and Scheduling, Games (Planning, Reasoning, Understanding, ...), Robotics (Autonomous driving, ...), Natural Language Processing (Text Generation, Understanding, Translation, Speech recognition, ...), Scientific Discovery (AI find new quasars, proved theorems, ...), Computer Vision (Autonomous vehicles, classification, segmentation, ...), ...

→ Current AI often isolated to single problems and not superhuman in every task

→ AI models are not without bias → fundamental differences in how AI perceives the world / environment

↳ can be tricked

↳ because of data it is fed / learning with

## ② AI Systems

**AI System** → can be defined as the study of rational agents and their environments

1. **Environment** „surroundings or conditions in which a person, animal or plant lives or operates“

↳ context of AI → surroundings of an (AI) agent and where agent operates

↳ does not have to be real. It could also be artificial.

Examples: Selfdriving cars : Streets, traffic, weather, road signs, pedestrians, ... ; Chess: chess board, chess pieces, ... ; ...



**Characteristics**: Discrete vs. Continuous → limited/countable number of distinct, clearly defined states → discrete, otherwise continuous (e.g. selfdriving car)

Observable vs. Partially Observable or Unobservable → possible to determine complete state of environment at each time point → observable, otherwise partially observable or unobservable

Difference: ...

Static vs. Dynamic → environment does not change while agent is acting → static (e.g. crossword puzzle), otherwise dynamic (e.g. You driving)

Single Agent vs. Multiple Agents → may contain other agents of same or different kind (agent design problems in multi-agent environment different from single agent)

Accessible vs. Inaccessible → agent can obtain complete and accurate information about state's environment → accessible environment, otherwise inaccessible

(empty room with state defined by temperature → accessible ; event on earth → inaccessible)

Deterministic vs. Non-deterministic/Stochastic → next stage completely determined by current state and actions of agent → deterministic, otherwise non-deterministic (stochastic event is random in nature → not determined by agent)

Episodic vs. Non-episodic/Sequential → episodic → each episode consists of agent perceiving and then acting. Quality of action depends just on an episode.  
sequential → agent requires memory of past actions

Importance: Different environments need different agents

Not every algorithm works in every environment (e.g. some algorithms do not work with uncertainty i.e. need perfect information)

Eg:	Environment	Discrete?	Observable?	Static?	Single Agent?	Accessible?	Deterministic?	Episodic?
Chess	Discrete	Observable	Static	Multi-Agent	Accessible	Deterministic	Sequential	

**Easy and Difficult Environments:**

Scenario 1: Static (no attention to possible environment changes), Observable (observe initial state at least), Discrete (possible actions can be enumerated) (easy)  
Deterministic (expected outcome of action is always identical to true outcome)

Scenario 2: Dynamic (changing even without acting), Partially Observable (current state?), Continuous (many possibilities), Stochastically (some actions in same state leads to different results) (difficult)

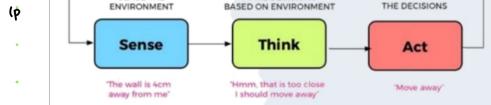
Questions to specify/define:

- infinite number of states?, - observe game perfectly?, - states/actions we cannot observe?, - environment changing?, - actions irreversible?,

- contain other agents?, - other agents have other forms than us?, - Opponent?, - observe complete environment?, - observation accurate?,

- actions always result in same next state?, - randomness in environment?, - determine endstate taking specific actions?, - memory of past needed to determine best action?

## 2. Agent



- Rules:
1. Agent must be able to perceive the environment
  2. The environments observations must be used to make decisions
  3. The decisions should result in action
  4. (The action taken by the agent must be rational)

Example of agents: Humans (senses, head, acts with body), Robots (sensors, decisions based on inputs, rules or a program, acts with motors for actuators), ...

→ A **rational agent** is an agent that does the "right thing"

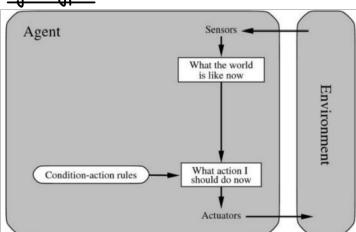
↳ rational agents are not **omniscient** → omniscient agent would know the actual outcome of its actions

↳ rational agents are not perfect → Rationality maximises expected performance → may not be optimal outcome (e.g. playing lottery has a negative expected outcome, better not to play, lost,...)

Measure performance: function that evaluates sequence of actions, task-dependent (e.g. vacuum different performance criteria (noise, time, energy,...) than a self-driving car (safety, comfort, routing,...))

- Design performance measure based on desired outcome, not desired agent behaviour (general rule)

### Agent Types:



**Reflex Agent:** Select action on basis of only current percept but ignores percept history.

→ Implemented through condition-action rules i.e. "map state to action"

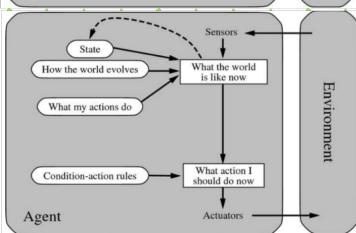
→ Makes very bad chess player → Does not look on the board, only at last move & no explicit goal of winning the game

### Problem

→ very limited in its decision making

→ No knowledge about anything which the agent cannot actively perceive

→ Can become very hard to handle (store, update,...) in complex environments



**Model-based Agent:** chose actions like reflex agents, but have a better comprehensive view of environment, i.e. keep track of world state

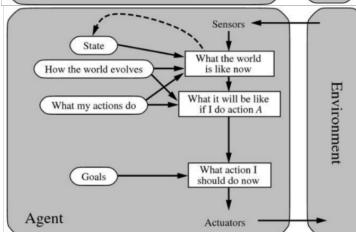
→ Input is not only interpreted, but mapped into internal state (and state dynamics) description (a world model)

→ Makes better chess agent → could keep track of current board situation when its percepts are only the moves

### Problem

→ How do actions affect the world?

→ What world model do I desire?



**Goal-based Agent:** Build on the information that a model-based agent stores but in addition knows desirable states

→ Agent knows desirable states and will try to choose actions accordingly

→ Main difference to previous approaches: takes decision-making into account (e.g. "What will happen if I do...?", "What will make me happy?")

### Problem

→ Things become difficult when long sequences of actions are required to find / achieve a goal



**Utility-based Agent:** Instead of providing goals, Utility-based agents use utility function providing a way to rate each action / state based on desired result

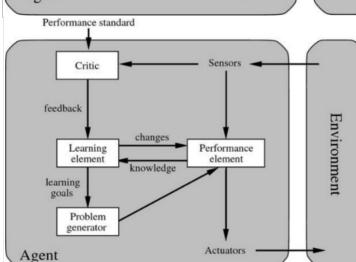
→ Goals provide binary distinction, while utility function provides continuous scale

→ Can help selecting between conflicting goals (e.g. speed or safety more important)

→ Certain goals can be reached in different ways (some ways are quicker, safer, more reliable,... (higher utility))

### Utility function

→ Maps a state or sequence of state onto a real number



**Learning Agents:** employ an additional learning element to gradually improve and become more knowledgeable over time about an environment

→ Can learn from its past experience, i.e. is able to adapt automatically based on experience

→ More robust toward unknown environments

Has 4 conceptual components:

1. **Learning Element:** responsible for making improvements by learning from environment

2. **Critic:** Gives feedback, describing how well agent is doing with respect to fixed measurement

3. **Performance Element:** responsible for selecting actions

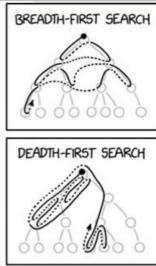
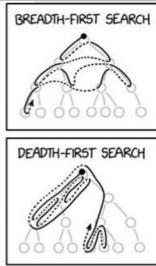
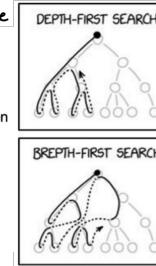
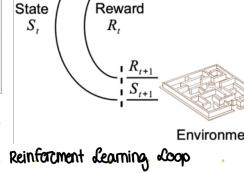
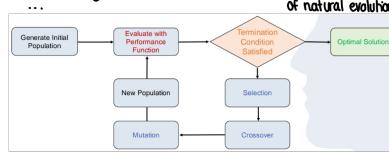
4. **Problem Generator:** suggesting actions which will lead to new experiences

Intelligent agents make intelligent actions → How do we decide what an intelligent action is and how do we select them?

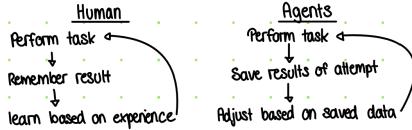
↳ Approaches: Search Algorithms: Understand/define "finding a good action" as search problem and use search algorithms; Most common search algorithms are tree-based ↳

Reinforcement learning: Developed within field of psychology, Trial and Error, Reactions/ACTIONS based on observation and experience

Genetic Algorithms: Survival of the fittest (Charles Darwin's theory of natural evolution)



Improve Performance:



### ③ Uninformed and Informed Search

- Problem-solving agents are result driven → Always focus on satisfying its goals, i.e. solve problem
- Problems given in human understandable way → reformulate problem for agent
- Problem-solving agents employ algorithms to develop / find solutions.

Steps to formulate solvable problem: formulate the goal → formulate problem given the goal

Components to formulate problem: (single state problem)

1. State Space and Initial State (Description of all possible states and initial environment)
2. Description of Actions (function that maps state to set of possible actions)
3. Goal Test (function to test whether current state fulfills goal definition)
4. Costs (cost function, mapping actions to costs; easy way is additive costs (sum of cost for all actions taken))

State Space: → set of all states reachable from initial state; defined by initial state and successor function → state-space graph

Path: sequence of states connected by sequence of actions

Solution: path from initial to goal state

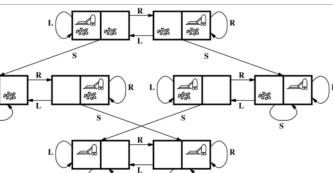
Optimal Solution: solution with min. path cost

E.g. → finding route on map → cost = gas use in relation to distance and time

↳ abstraction to gain problem understanding can be used to solve real world complexity

↳ 8-Puzzle:

States → location of tiles (ignore during sliding); Goal test → situation corresponds to goal state; Path cost → Number of steps in path; Actions → How blank tile (left, right, up, down)



↳ easier than separate moves for each tile  
↳ ignore actions like jumping if stuck etc.

↳ 8-Queens Problem: (placing 8 queens on an chessboard without threatening each other)

States → Any configuration of 8 queens on board; Goal test → no pair threatening each other; Path cost → Not of interest; Actions → move queen to other square

↳ inefficient complete-state formulation →  $64 \cdot 63 \cdots 57 \approx 3 \cdot 10^{14}$  states!

States → n non-attacking queens in left n columns; Goal test → //

; Path cost → //; Actions → add queen in column n+1, without

attacking others



Planning Problem: transform initial state into goal considering future actions and their outcomes

Search: Process of finding (optimal) solution to such problem in form of sequence of actions.

Tree Search Algorithms → treat state-space graph as tree → use simulated iterative exploration of state space; needs strategy to determine next node

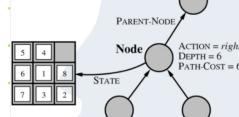
E.g. Route finding → initial state node = start city → next nodes cities reachable from initial city → and so on with any city

↳ fringe: set of all nodes at the end of all visited paths (frontier/border)

Depth: Number of levels in tree

state: Representation or physical configuration.

Node: Data structure to present part of search tree. Includes state, parent node, take action, path costs, current depth



↳ Expanding a node creates new nodes (leaf-successor nodes of chosen node) and updates tree

Search Strategy: defined by picking order of node expansion

Evaluated along: Completeness (does it find solution if one exists), Time Complexity (Number of Node expansions), Space Complexity (Max. number of nodes in memory), Optimality (does it always find the optimal solution)

Time and Space complexity measured in terms of:

- b: max. branching factor of search tree

- d: depth of optimal solution

- m: max. depth of state space tree

## Two main groups of search strategies:

### Uninformed Search

(have no information except problem definition)

→ Breadth-first search (BFS)

→ Uniform-cost search

→ Depth-first search (DFS)

→ Depth-limited search

→ Iterative deepening

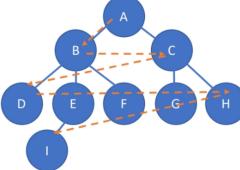
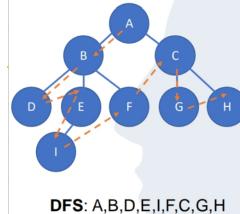
### Informed Search

(have additional knowledge about problem and an idea where to "look" for solutions)

→ Greedy Best-first Search

→ A\* Search

→ Memory-Bounded Heuristic Search



### Iterative Deepening Search

→ Increase  $l$  after each failed search

↳  $l = 1, 2, 3, \dots$

Completeness: Yes

Time Complexity: first level searched  $d$  times  $\rightarrow d \cdot b + (d-1)b^2 + \dots + 1 \cdot b^d = \sum_{i=1}^d (d-i+1) \cdot b^i$

Space Complexity: linear  $\rightarrow O(b \cdot d)$

Optimality: Yes, shortest path is found

### Bidirectional Search

→ two search simultaneously, starting with root and goal.

Stop if node occurs in both searches

→ Reduction in complexity ( $b^{d/2} + b^{d/2} < b^d$ )

→ Only possible if actions reversible

→ Search paths may not meet for depth-first bidirectional search

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes*	Yes*	No	Yes, if $l \geq d$	Yes
Time	$b^{d+1}$	$b^{[C^*/\epsilon]}$	$b^m$	$b^l$	$b^d$
Space	$b^{d+1}$	$b^{[C^*/\epsilon]}$	$bm$	$bl$	$bd$
Optimal?	Yes*	Yes	No	No	Yes*

### Uninformed Tree Search Strategies

#### Uniform-Cost Search (UCS)

→ Node associated with fixed cost and they accumulate over path within search. UCS uses lowest cumulative cost to find path

Breadth-First Search (BFS) → conceptually simple, but memory consumption can be too costly (depth 1 → 107 kB, depth 16 → 600 kB)

→ Special UCS case → all costs equal. Start at root and explore bidirectional by level. BFS stops as soon as goal generated → UCS explores all nodes at goal's depth → check whether one has lower cost.

Completeness: Yes, if each step has positive cost → otherwise, infinite loops possible (BFS complete as well)

Complexity: BFS →  $O(b^d)$ , UCS →  $O(b^{d+\lfloor \log_b(C/\epsilon) \rfloor})$  every action costs at least  $\epsilon$

Optimality: Yes, since nodes expand in increasing order of path costs → BFS optimal.

BFS & UCS often implemented using priority queue ordered by costs

→ exponential-complexity search cannot be solved by uninformed methods for any but smallest instances

#### Depth-First Search (DFS)

→ Starts at tree root and explores tree as far as possible along one branch.

Then going step-wise back and explore alternative branches.

Completeness: No, fails in infinite-depth search spaces and loops

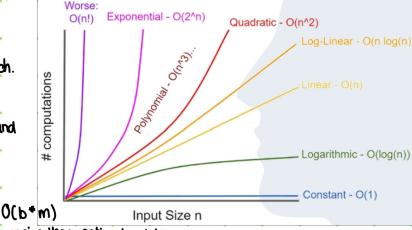
↳ can be modified to complete by avoiding repeated states and limit depth

Time Complexity:  $O(b^m)$  → each branch until max. depth

↳ terrible if  $m > d$ , good in dense settings

Space Complexity: Only branch and unexpanded siblings stored → linear  $O(b \cdot m)$

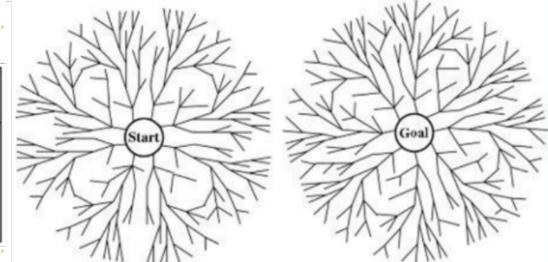
Optimality: No, longer solutions may be found before shorter → more expensive than optimal solution



Criteria	BFS	DFS
Concept	Traversing tree level by level	Traversing tree sub-tree by sub-tree
Data Structure (Queue)	First In First Out (FIFO)	Last In First Out (LIFO)
Time Complexity	$O(\text{Vertices} + \text{Edges})$	$O(\text{Vertices} + \text{Edges})$
Backtracking	No	Yes
Memory	Requires more memory	Less nodes are stored normally (less memory)
Optimality	Yes	Not without modification
Speed	In most cases slower compared to DFS	In most cases faster compared to BFS
When to use	If the target is relatively close to the root node	If the goal state is relatively deep in the tree

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
    RECURSIVE-DLS(MAKE-NODE([INITIAL-STATE[problem]]), problem, limit)
function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
    cutoff-occurred? ← false
    if GOAL-TEST(problem, STATE[node]) then return node
    else if DEPTH[node] = limit then return cutoff
    else for each successor in EXPAND(node, problem) do
        result ← RECURSIVE-DLS(successor, problem, limit)
        if result = cutoff then cutoff-occurred? ← true
        else if result ≠ failure then return result
    if cutoff-occurred? then return cutoff else return failure
```

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution
inputs: problem, a problem
for depth ← 0 to ∞ do
    result ← DEPTH-LIMITED-SEARCH(problem, depth)
    if result ≠ cutoff then return result
end
```



## Informed Tree Search Strategies:

→ improve performance; giving hints about desirability of different states  
 Uninformed search algorithms are inefficient. → Bring knowledge into process of what node to expand to → e.g. straight-line distances may be good approximation for remaining travel distance

Heuristics h: → can go wrong!

→ Informally denotes a "rule of thumb", i.e. rule that may be helpful in solving the problem. In tree-search, heuristic denotes function  $h \rightarrow$  estimates remaining costs to reach goal.

### Greedy Best-first Search:

→ Using evaluation function  $f(n) = h(n)$  to estimate cost from node n to goal

↳ e.g.  $h(n) = \text{hLD}(n)$  = straight-line distance from n to Bucharest (on map)

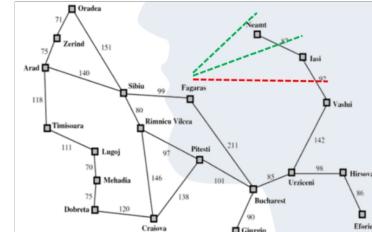
→ Expand node with smallest cost according to  $f(n)$ , i.e. here the heuristic

(Completeness: NO (loops) → complete in infinite state space (no loops))

Time Complexity: Worst case  $O(b^m)$  → same as DFS → improved using good heuristics

Space Complexity: Worst case  $O(b^m)$  → all nodes in memory

Optimality: No, solution depends on heuristic



### Example:

→ heuristic not helping

From Iasi to Fagaras

Problem → meant "closer" →

Fagaras than Vaslui

### A\* Search:

→ Build on best-first search → tries to minimize estimated  $h(n)$  and true  $g(n)$  cost.

Idea: Avoid expanding paths that are already expensive → evaluate complete not only remaining cost

→ i.e.  $f(n) = g(n) + h(n)$

Evaluation function:  $f(n)$  → cost so far to reach n

$h(n)$  → estimated cost to get from n to goal

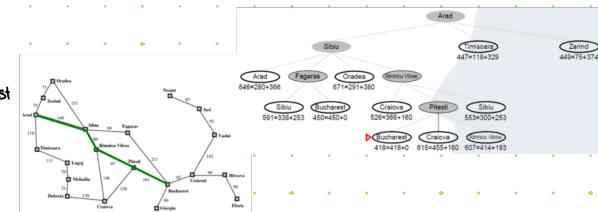
$f(n)$  → estimated cost of path to goal via n

Completeness: Yes, except if there are infinite nodes with  $f(n) \leq f(g)$

Time Complexity: number of nodes grow exponentially unless error of  $h(n)$  is bounded by logarithm of value of actual path cost  $h^*(n)$ , i.e.  $|h(n) - h^*(n)| \leq O(\log h^*(n))$

Space Complexity: keep all nodes in memory

Optimality: Depends on heuristic



d	Search Cost (nodes generated)			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	245	1.79	1.79
4	112	13	12	2,87	1.48	1.45
6	680	20	18	2,73	1.34	1.30
8	6384	39	30	2,80	1.33	1.24
10	47127	93	30	2,79	1.38	1.22
12	3644035	227	73	2,78	1.42	1.24
14	—	539	113	—	1.44	1.23
16	—	1301	211	—	1.45	1.25
18	—	3025	38	—	1.46	1.26
20	—	7276	676	—	1.47	1.27
22	—	18094	1219	—	1.48	1.28
24	—	39135	1641	—	1.48	1.26

### Kingdom of Heuristics

Admissible Heuristics: → admissible if never overestimates cost to reach goal; formally:  $h(n) \leq h^*(n)$  →  $h^*(n)$  true cost from n to goal; e.g. straight-line distances ( $h_{SL}(n)$ ) admissible for actual distance

Consistent Heuristics: → consistent if for every n and every successor n' generated by action a it holds  $h(n) \leq c(n, a, n') + h(n')$ .

When going from neighboring nodes a to b, heuristic difference/step cost never overestimates actual step cost.

- ↳ if route from n to goal cheaper than  $h(n)$  → violates property that  $h(n)$  lower bound on cost to reach the goal
- Every consistent heuristic is admissible.
- If  $h(n)$  consistent, values of  $f(n)$  along any path are non-decreasing.

### Designing Heuristics: e.g. for 8-puzzle

Two ideas  $h(n) = 1\text{-misplaced tiles} \rightarrow h(\text{start}) = 8 \rightarrow h_{1\text{-mis}}(n)$

( $h_{1\text{-mis}}(n)$  = Total Manhattan distance (squares from desired tile location) →  $h_2(\text{start}) = 3+1+2+2+2+2+3+3+2 = 18$ )

$h(\text{start}) = 26$

### Heuristics for Relaxed Problems:

Relaxed Problems: Problem with fewer restrictions on actions

The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem.

e.g. If rules of 8-puzzle relaxed → tile can move anywhere

→ HITS gives shortest solution

If rules are relaxed → tile can move to any adjacent square

→ HTAN gives shortest solution

→ Looking for relaxed problems is a good strategy for

inventing admissible heuristics.

### Optimality of Heuristics:

If  $h(n)$  is admissible,  $A^*$  using tree-search is optimal.

Proof:



$C^* = g(G) + h^*(G)$

$f(G) = C^* < f(G'')$

$h(G) \leq h^*(G)$

because  $h(G) = 0$  (holds for all goal nodes)



$f(G) = C^* < f(G'')$

$h(G) = h^*(G)$

because  $h(G) = 0$  (holds for all goal nodes)

Assumption: If  $A^*$  selects path p, p is shortest path.

Proof: Assume contradiction → p' is shortest.

• Moment before p is chosen from fringe.

• Some part of p' will also be on fringe → partial path p''.

• p was expanded before p'',  $f_p(p) \leq f_{p''}(p'')$

• p is goal,  $h(p) = 0$ ,  $cost(p) \leq cost(p'') + h(p'')$

• h admissible →  $cost(p'') + h(p'') \leq cost(p) + h(p)$  for any p' to goal extending p'

• Thus  $cost(p) \leq cost(p'')$  for any other path p' to goal → contradicts assumption that p' is shortest.

### 3 Alternatives to A\*: Memory-Bounded Heuristic Search:

Problem: A\* (B\*; best-first) space problem

#### 1. Iterative Deepening A\* (IDA\*)

- like iterative deepening → cutoff information is f-cost (g+h) instead of depth
- Recursive best-first search (RBFS)
- Recursive algorithm → attempts mimicking standard best-first search with linear space
- keeps track of f-value of best alternative
- path available from any ancestor of current node and heuristic evaluations are updated with result of successors

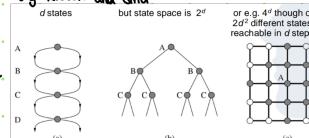
#### 3. (Simple) Memory-bounded A\* (CSA\*)

- drop worst leaf node when memory full
- its value updated to its parent
- may need to be re-searched later

### Graph Search:

→ failure detecting repeated states can turn linear to exponential problem

#### E.g. Ribbon and Grid



→ Remember visited states in list → e.g. Dijkstra's algorithm is graph-search variant of uniform cost search

```
function TREE-SEARCH(problem) returns a solution, or failure
    initialize the frontier using the initial state of problem
    loop do
        if the frontier is empty then return failure
        choose a leaf node and remove it from the frontier
        if the node contains a goal state then return the corresponding solution
        expand the chosen node, adding the resulting nodes to the frontier
```

```
function GRAPH-SEARCH(problem) returns a solution, or failure
    initialize the frontier using the initial state of problem
    initialize the explored set to be empty
```

```
loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
    only if not in the frontier or explored set
```

### Tree Search vs. Graph Search:

Tree Search: node represents possible path to associated domain state

Problem: later detected path may be better → previous found paths reinvestigated with new, cheaper path

Solutions: 1. Add ability to detect repeated states → graph search → works also for BFS, DFS, ...

2. Ensure cheaper path is always taken first → consistent heuristics → specific solution for A\*

## 4 Local Search and Adversarial Search

→ larger spaces, path isn't goal

Optimization Problem: → all states/nodes can be solution (to different degrees) → target is finding state which optimizes (min, max, ...) solution according to objective function → finding "best" states  
→ no (explicit) goal state and no path (costs).

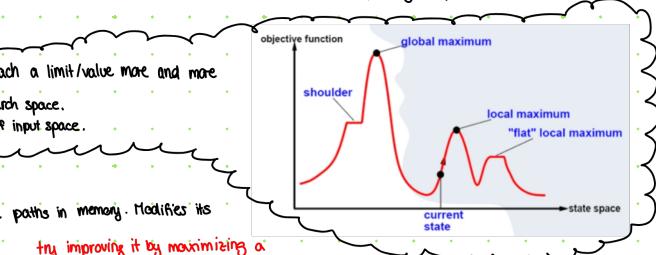
Objective/Evaluation Function: → tells how good a state is (also compared to other states). Value either minimized or maximized depending on optimization problem.

e.g. Darwinian evolution and survival of the fittest → optimization process

Convergence: Describes property of sequence of (function) values to approach a limit/value more and more

Global Optimum: Extremum (min. o. max.) of objective function for entire input search space.

Local Optimum: " " " " " for given region of input space.



### Local Search Algorithms (Iterative Improvement Methods):

Local Search: traversing only a single state rather than saving multiple paths in memory. Modifies its state iteratively, trying to improve a criteria.

Goal is to find state fulfilling all goal constraints not path.

Advantages: → little/constant amount of memory usage  
→ find reasonable solution in very large state spaces

Disadvantages: → No guarantees for completeness or optimality.

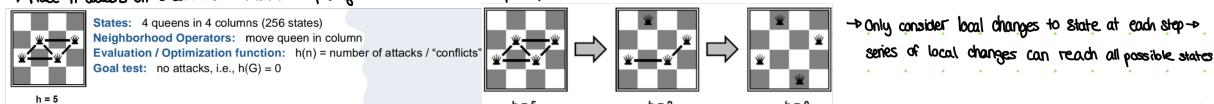
#### Example:

##### Travelling Salesman Problem

→ start with complete (prob. suboptimal) tour → perform pairwise exchanges → variants of approach get within 1% of optimal solution quickly with thousands of cities ↗

##### n-Queens Problem

→ Place n-Queens on chessboard without capturing each other → Move queen, so reduces number of conflicts



### Hill Climbing / Greedy Local Search:

→ expand current state → move to neighbour with highest evaluation → do so until maximum reached  
(evaluation goes down again)

Problem: Local Optima → algorithm stops at hilltop → hill does not have to be optimal

Idea: Randomized Hill Climbing Variants

1. Randomized Restart Hill Climbing → different initial positions result in different local optima.
2. Stochastic Hill-climbing → select successor node randomly → better nodes have higher probability of being selected

#### Ridge Problem / Issue:

→ Every neighbour state appears to be downhill → search space has an uphill → neighbours not states / steps (discrete)

#### Example: 8-Queens Problem

→ Heuristic h: → number of pairs of queens that attack each other, if we move a queen in its column to that square

$h=17 \rightarrow$  Best neighbour(s):  $h=12$

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
```

```
inputs: problem, a problem
local variables: current, a node
neighbor, a node
```

```
current ← MAKE-NODE(INITIAL-STATE[problem])
```

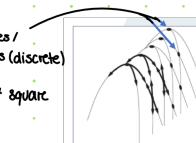
```
loop do
```

```
    neighbor ← a highest-valued successor of current
```

```
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
```

```
    current ← neighbor
```

```
end
```

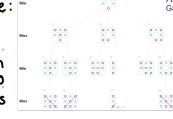




Game Trees: → to solve games

- ↳ arranged in levels corresponding to players
- ↳ root → active / current player
- leaf nodes called terminal → each terminal node has utility value corresponding to game outcome.

Tic-tac-toe:



→ small tree

→ fewer than  $9! = 362,880$

terminal nodes

→ in chess →  $10^{40}$

### Alpha-Beta Pruning:

→ modified, optimized Minimax version

→ pruning to reduce amount of exploration without losing correctness of minimax

Based on two parameters:

Alpha: best (highest-value) choice found so far at any point along path of maximizer to root. Initial value of alpha → -∞.

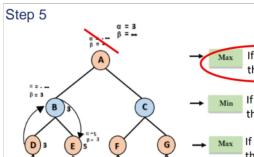
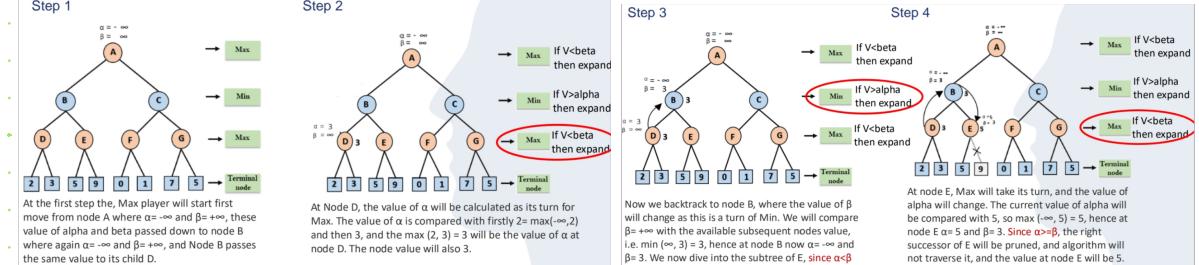
Beta: best (lowest-value) choice found so far at any point along path of minimizer to root. Initial value of beta → +∞.

→ Remove all nodes not really affecting final decision but slowing algorithm, hence by pruning these nodes.

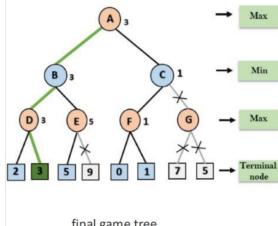
Key difference to Minimax: Max player only update alpha value. Min player update beta value.

→ While backtracking tree, node values passed to upper nodes instead of values of alpha and beta. → Only pass alpha, beta values to child nodes.

→ Video explanation link on slide 55.



At the first step, the Max player will start move from node A where  $\alpha = -\infty$  and  $\beta = +\infty$ , these value of alpha and beta passed down to node B where again  $\alpha = -\infty$  and  $\beta = +\infty$ , and node B passes the same value to its child D.



slide 62 → outlook AlphaGo video

A lot of AI tasks are intractable → Search is "only" way to handle them

Many applications require some form of search (e.g. Learning, Reasoning, Planning, Natural Language Understanding, Computer Vision, ...)

↳ Even in current State-of-the-Art algorithm, search has still a key role in their success

A lot of problems can be abstracted or displayed as a search problem!

Teaching AI to play games:

Games...

1. are useful metric (allows improved tracking, quantifiable)
2. provide safe place to train and evaluate (low-stake environments, safer training than real world)
3. captures interest and imagination (showcase advancements, opens advancements to public (AlphaGo vs. Lee Sedol))

Games vs. Search Problems?

1. "Unpredictable" opponent (specify move for every possible reply, different agents)
2. time limits (unlike math, how long to find goal, need approximation)
3. Most games are deterministic, turn-taking, two-player, zero-sum
4. But most "real" problems are stochastic, parallel, multi-agent, utility based

Problem of massive game trees: → tree too big to fully traverse.

(e.g. chess ( $b \approx 35$ ) tree grows like  $1, 3, 32, 325, 4225, 45000, \dots$  nodes.)

In Go  $b \approx 250 \rightarrow$  more possible Go positions than atoms in universe!)

Heuristic value of functions → limit depth → for each leaf calculate value for current situation based on heuristics

Pruning → cut back tree → ignoring unwanted portions which make no difference to final result (e.g. Alpha-Beta Pruning) (or randomly select subtree → stochastic search)

Estimation of likely outcome if current situation based on heuristics continue

Minimizer // initial value of beta → +∞.

Maximizer // initial value of alpha → -∞.

→ Remove all nodes not really affecting final decision but slowing algorithm, hence by pruning these nodes.

Key difference to Minimax: Max player only update alpha value. Min player update beta value.

→ While backtracking tree, node values passed to upper nodes instead of values of alpha and beta. → Only pass alpha, beta values to child nodes.

Minimax Algorithm:

→ game tree → nodes represent states of game → edges " moves by players

player are:

- MIN: decrease chances of MAX to win game. (Opponent)

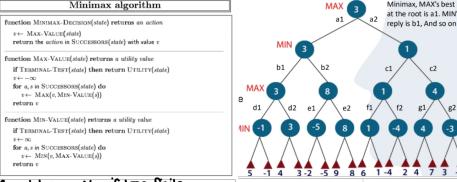
- MAX: increases chances of winning

Game played turn by turn following above strategy

→ choose move position with highest minimax value → assuming opponent plays best response to action

Yellow DFS concept → each level either MIN o. MAX → you MAX & opponent MIN

→ On MAX level we maximize value → on MIN level we minimize



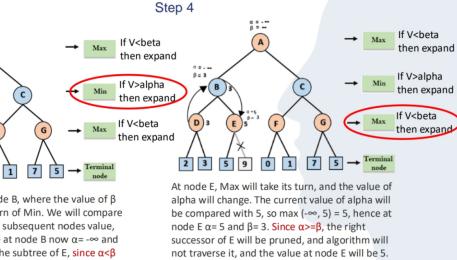
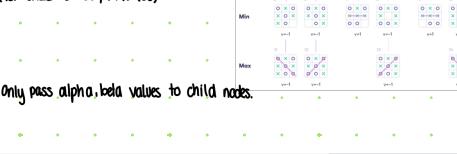
Completeness: Yes, if tree finite

Time Complexity:  $O(b^m)$

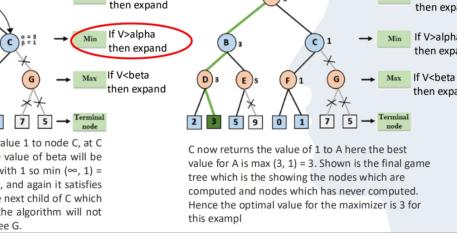
Space Complexity:  $O(b \cdot m)$

Optimal: Yes, assuming optimal opponent

(for chess  $b \approx 35$ ,  $m \approx 100$ )



Now we backtrack to node B, where the value of  $\beta$  will change as this is a turn of Min. We will compare  $\beta = +\infty$  with the available subsequent nodes value, i.e.  $\min(\infty, 3) = 3$ ; hence at node B  $\alpha = -\infty$  and  $\beta = 3$ . We now dive into the subtree of E, since  $\alpha < \beta$



Node E will take its turn, and the value of  $\alpha$  will change. The current value of  $\alpha$  will be compared with 5, so  $\max(-\infty, 5) = 5$ , hence at node E  $\alpha = 5$  and  $\beta = 3$ . Since  $\alpha > \beta$ , the right successor of E will be pruned, and the algorithm will not traverse it, and the value at node E will be 5.

Now we backtrack to node B, where the value of  $\beta$  will change as this is a turn of Min. We will compare  $\beta = 3$  with left child which is 0, and  $\max(0, 3) = 3$ , and then compared with right child which is 1, and  $\max(3, 1) = 3$ ; still  $\alpha$  remains 3, but the node value of F will become 1.

→ can save a lot of calculations, if tree is well ordered

→ in ideal ordering (best move first) reduce complexity to  $O(b^{m/2})$  instead of  $O(b^m)$

→ to optimize Alpha-Beta Pruning, you have to predict interesting path to increase amount of cutoffs

Problems:

→ needs fast evaluation function

→ In games with large branching factors (e.g. Go) exploration of Alpha-Beta not good enough (to explore full minimax tree depth 3 →  $300^3 = 27.000.000$  evaluations)

AlphaGo and AlphaZero:

→ In 2016, Deepmind introduced AlphaGo → combination of (Monte-Carlo) Tree Search and Neural Nets to beat Go

A lot of AI tasks are intractable → Search is "only" way to handle them

Many applications require some form of search (e.g. Learning, Reasoning, Planning, Natural Language Understanding, Computer Vision, ...)

↳ Even in current State-of-the-Art algorithm, search has still a key role in their success

A lot of problems can be abstracted or displayed as a search problem!

# 5 Constraint Satisfaction Problems

→ technique where problem is solved when solution satisfies certain constraints OR rules of problem

Components: - state, defined by variables  $X_i$  with values from domain  $D_i$

- Goal test, defined as set of constraints  $C$  specifying allowable combinations of values for subsets of variables

Solving CSP requires: - state-space, notion of solution

Real world CSPs: Assignment problems, Timetabling problems, Hardware configuration, Spreadsheets, Scheduling, Floor planning, ...

↳ involve real-valued variables → linear constraints solvable in polynomial time → linear programming

→ nonlinear constraints undecidable

Example: n-Queens → place n queens on n × n chessboard → no two queens threaten each other

Variables:  $X =$  one variable for each row; Domain of variables:  $D =$  number between 1 and 8; Constraints:  $C =$  Enumeration of disallowed configurations

Constraint Graphs: → Abstraction of problem → easier to solve and understand

→ Every variable represented as node

→ Every edge indicates constraint between them

Example: Four Color Map Theorem



Problem: assign each territory a color such that no two adjacent territories have same color

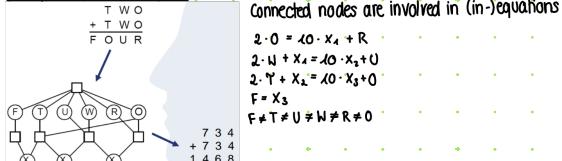
Variables:  $X = \{WA, NT, Q, NSW, V, SA, T\}$

Domain of variables:  $D = \{r, g, b, y\}$

Constraints:  $C = \{SA \neq WA, SA \neq NT, SA \neq Q, \dots\}$

→ Always possible to color map using 4 or less colors → ensuring regions sharing a border are different colors (videos on slides 13 & 14)

Example: TWO + TWO = FOUR



Example: Wildlife Corridor Design → slide 16

Search:

Map CSPs into search problems: Nodes = assignments of values to subset of variables

Neighbors of a node = nodes in which values are assigned to one additional variable  
Start node = empty assignment

Goal node = node which assigns value to each variable and satisfies all constraints

Build simple search space:

Start node = empty assignment of values to variables  
its successors are all possible ways of assigning a value to a variable ( $d=1$ )  
their successors are those with two variables assigned ( $d=2$ )  
...

Commutative Variable Assignments:

→ Not relevant if we do [WA = red then NT = green] or [NT = green then WA = red], assignments are commutative, order not important

→ At each node, we only need to make assignments for one of the variables → reduces total complexity to  $v^n$



Example Sudoku: → CSP with 81 variables

Constraints: → some values are assigned in the start

→ 27 constraints on 9 values that must all be different (9 rows, 9 columns, 9 squares)

Constraint Propagation: → list of possible values in empty fields  
→ successively eliminate values

Status: Automated constraint solvers can solve hardest puzzle in no time

Assignment of values to variables:

- state in state-space not "blackbox" anymore → defined by assigning values to some or all variables such as  $X_1 = U_1, X_2 = U_2, \dots$

Can be done:

1. Consistent/legal assignment → does not violate any constraint or rules

2. Complete assignment → every variable is assigned with value and solution to CSP remains consistent

3. Partial assignment → assigns values only to some variables

→ don't care about path to solution → solutions described by constraints → can do better than search trees → still maintain data structure (compared to local search)

Types of constraints:

Unary constraints: involve a single variable (e.g. South Australia ≠ green)

Binary constraints: involve pairs of variables (e.g. South Australia ≠ Western Australia)

Higher-order constraints: involve 3 or more variables (e.g.  $2 \cdot W + X_1 = 10 \cdot X_3 + U$ )

Preferences (soft constraints): e.g. red better than green; not binding → respect as many as possible  
→ constrained optimisation problems

Two principal approaches to solve CSPs:

1. Search: → successively assign values to variable → check all constraints → if constraint violated → backtracking → until all variables have assigned values
2. Constraint Propagation: → maintain set of possible values  $D_i$  for each variable  $X_i$  → try to reduce size of  $D_i$  by identifying values that violate some constraints

Naïve Search: → slide 20 video

Complexity:  $n$  variables → all solutions are at depth  $n$  in search tree  
all variables have  $v$  possible values

↳ at level 1:  $n \cdot v$  possible assignments → choose one of  $n$  variables and one of  $v$  values for it

↳ at level 2:  $(n-1) \cdot v$  possible assignments for each previously unassigned variable → choose one of remaining  $n-1$  variables and one of  $v$  values for it

In general: branching factor at depth 1:  $(n-1) \cdot v$

→  $n \cdot (n-1) \cdot v \cdot v$  leaves in total at  $d=2$

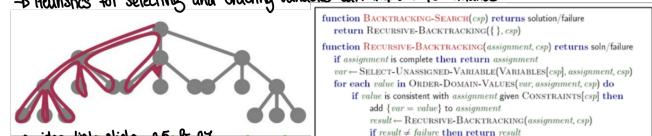
→ Therefore, search tree has  $n! v^n$  leaves

Backtracking Search → DFS with single variable assignment per level

↳ basic uninformed search algorithm for CSPs (add one constraint at a time without conflict, succeed if legal assignment is found → can solve n-queens for up to  $n = 25$ )

→ Complexity in worst case still exponential

→ Heuristics for selecting and ordering variables can improve performance

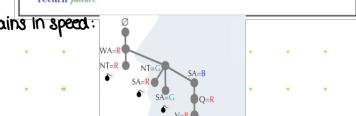


General-purpose methods can give huge gains in speed:

1. Next variable to design?
2. Order in which its value is tried?
3. Detect inevitable failures early?
4. Take advantage of problem structure?

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
    return RECURSIVE-BACKTRACKING(1, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
    if assignment is complete then return assignment
    else ← SELECT-UNASSIGNED-VARIABLE(csp, assignment, csp)
        for each value in ORDER-DOMAIN-VALUES(csp, assignment, csp) do
            if value is consistent with assignment given CONSTRAINTS(csp) then
                add {var = value} to assignment
                result ← RECURSIVE-BACKTRACKING(assignment, csp)
                if result ≠ failure then return result
                remove {var = value} from assignment
        return failure
```



## Heuristics for CSPs:

Domain-Specific Heuristics → depend on particular characteristics of problem

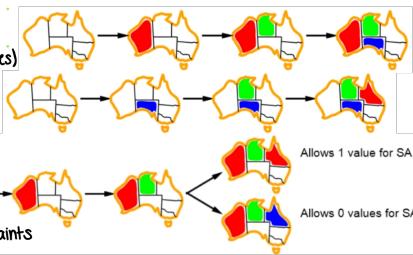
General-purpose Heuristics → for CSP good general-purpose heuristics:

↳ Minimum Remaining Values Heuristic (choose variable with fewest consistent values)

↳ Degree Heuristic (choose variable with most constraints on remaining variables)

↳ Least Constraining Value Heuristic (variable given, choose value that rules out fewest values in remaining variables)

→ Used in this order, these can greatly improve search speed



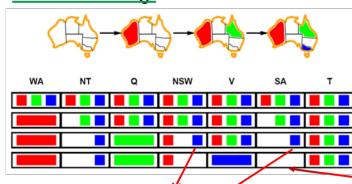
## Node Consistency:

→ variable (node) is consistent if possible values of this variable are conform to all unary constraints

↳ Example: Sudoku → some nodes already filled out i.e. constrained to a single value

Local consistency → defined by a graph where each of its nodes is consistent with its neighbors. Done by iteratively enforcing the constraints corresponding to edges.

Forward Checking: → keep track of remaining legal values for unassigned variables and terminate search when any variable has no more legal values



Arc Consistency: → improve performance by looking at larger sets of constraints and how they interact

Arc: constraint involving two variables is called arc or binary constraint

Arc Consistency: arc consistent if for each value of  $X$  in domain of  $X$  there is value  $Y$  in domain of  $Y$  such that constraint arc  $(X, Y)$  is satisfied.

$\forall X \in \text{dom}(X), \exists Y \in \text{dom}(Y)$  such that arc  $(X, Y)$  is satisfied

→ can be generalized to  $n$ -ary constraints → each tuple involving  $X_i$  has to be consistent

function AC-3( $csp$ ) returns the  $CSP$ , possibly with reduced domains

inputs:  $csp$ , a binary  $CSP$  with variables  $\{X_1, X_2, \dots, X_n\}$

local variables:  $queue$ , a queue of arcs, initially all the arcs in  $csp$

while  $queue$  is not empty do

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(queue)$

if REMOVE-INCOSISTENT-VALUES( $X_i, X_j$ ) then

for each  $X_k \in \text{NEIGHBORS}[X_i]$  do

add  $(X_i, X_k)$  to  $queue$

If  $X$  loses a value, neighbors of  $X$  need to be rechecked.

function REMOVE-INCOSISTENT-VALUES( $X_i, X_j$ ) returns true iff succeeds

removed ← false

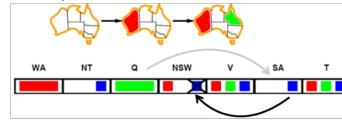
for each  $X_k \in \text{DOMAIN}[X_j]$  do

if no value  $y$  in  $\text{DOMAIN}[X_j]$  allows  $(x, y)$  to satisfy the constraint  $X_i \rightarrow X_j$

then delete  $x$  from  $\text{DOMAIN}[X_j]$ ; removed ← true

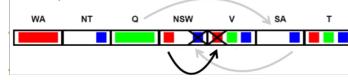
return removed

AC-3 algorithm



Maintaining Arc Consistency (MAC): → after each new assignment of a value to a variable, possible values of neighbours have to be updated:

If one variable (NSW) loses value (blue), need to recheck its neighbours as well because they might have lost a possible value.



Path Consistency: Arc Consistency often sufficient to solve the problem (all variable domains reduced to size 1) or show that problem cannot be solved (some domains empty)

↳ not enough → always consistent value in neighboring region

↳ tightens binary constraints by considering triples of values

pair of variables  $(X_i, X_j)$  path-consistent with  $X_m$  if for every assignment that satisfies constraint on arc  $(X_i, X_j)$  there is assignment that satisfies the constraints on the arcs  $(X_i, X_m)$  and  $(X_j, X_m)$

e.g. Sudoku: simple puzzles solved with AC-3 → 27 constraints → each requires that 9 values are all different

→ These 9-valued ANDIF constraints can be converted into pairwise binary constraints ( $9 \times 8/2 = 36$ ) →  $27 \cdot 36 = 972$  arc constraints

→ not all problems can be solved with constraint propagation alone → search

$k$ -Consistency: concept generalized → set of  $k$  values need to be consistent: 1-consistency = node consistency ○

2-consistency = arc consistency ○○

3-consistency = path consistency ○○○

→ May lead to faster solution but checking for  $k$ -consistency is exponential in  $k$  in the worst case

→ therefore arc consistency more frequently used in practice

## Constraint Propagation and Backtracking Search:

→ each time a variable is assigned, constraint propagation algorithm run to reduce number of choice points in search

↳ improve speed of backtracking search further

Possible Algorithm: Forward Checking, AC-3 → initial queue of constraints only contains constraints with variable that has been changed

Example videos: slides 45 + 46

## Local Search for CSPs:

Modifications for CSPs: work with complete states, allow states with unsatisfied constraints, operators reassigned variable values

Hill-conflicts Heuristic: random select conflicted variable, choose value that violates fewest constraints, hill-climbing with  $h(n) = \#$  of violated constraints

Performance: can solve randomly generated CSPs with a high probability, except in a narrow range  $R$

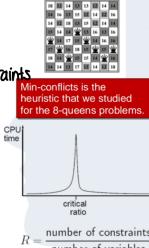
## Power of Problem Decomposition:

Assumption: Search space for a constraint satisfaction with  $n$  variables, each of which can have  $d$  values =  $O(d^n)$

Idea: Decomposing problem into subproblems with  $c$  variables each → each problem has complexity =  $O(d^c)$ , there are  $n/c$  such problems →

Total complexity =  $O(n/c \cdot d^c)$  → unconditional independence is rare!

Effect: total complexity can be reduced from exponential in  $n$  to linear in  $n$  (assuming  $c$  is constant) e.g.  $n=80, d=2, c=20 \rightarrow 2^{80} \neq 4 \cdot 2^{20}$



tree-structured CSPs → tree-structured if in the constraint graph any two variables are connected by a single path



Theorem: any tree-structured CSP can be solved in linear time in number of variables ( $O(n \cdot d^2)$ )

### Linear Algorithm:

1. choose variable as root → order nodes so that parent always comes before its children (each child only one parent)

2. For  $j = n$  down to 2 → Make arc( $X_i : X_j$ ) arc-consistent, calling Remove-Inconsistent-Value( $X_i, X_j$ )

3. For  $i = 1$  to  $n$  → Assign to  $X_i$  any value consistent with its parent

Problem: tree-structured problems are rare → approaches for making problems tree-structured:

1. Cutset Conditioning → removing nodes so that remaining nodes form tree

2. Collapsing nodes together → Decompose graph into set of independent tree-shaped subproblems

### Cutset Conditioning:

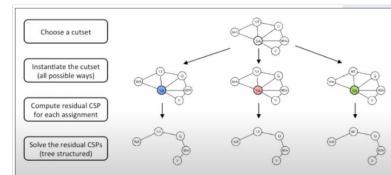
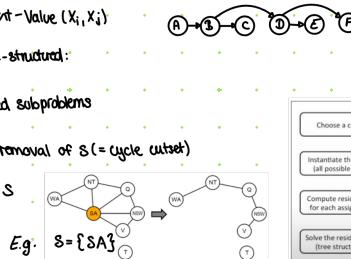
1. Choose subset  $S$  of variables such that constraint graph becomes tree after removal of  $S$  (= cycle, cutset)

2. Choose (consistent) assignment of variables for  $S$

3. Remove from remaining variables all values inconsistent with variables of  $S$

4. Solve CSP problem for remaining variables

5. If no solution → choose different assignment for variables in  $S$



## 6 Logic and AI 1: Propositional Logic

→ make agent smarter with logical reasoning

Syntax → sequence of specific language which should be followed in order to form a sentence. Representation of language & related to grammar and structure.

Semantics → sentence or syntax which logic follows should be meaningful. Defines sense of sentence which relates to real world.

e.g. Programming → Syntax: structural rules, i.e. separate statements with semi-colon, ... (syntax correct if compiled); Semantics: Does program work? Makes sense?

Logic → key behind any (formal) knowledge. Allows person to filter necessary information and draw conclusion. In AI representation of knowledge done via logics.

Knowledge Base (KB) → represents actual facts existing in real world. Central component of knowledge-based agent. Set of sentences describing information related to the world. → set of sentences in formal language.

Inference engine	domain-independent algorithms	declarative approach to building an agent (tell it what it needs to know)
Knowledge base	domain-specific content	• ask itself what to do → answers should follow from KB • Agents can be viewed at knowledge level i.e. what they know

### Agents:

Intelligent Agents: goal-directed agent. Perceives its environment through its sensors using observations and built-in knowledge → acts upon environment through its actuators

Knowledge-based Agents: Represent states, actions, etc...; Incorporate new percepts; update internal representations of world; deduce hidden properties of world & appropriate actions

```
function KB-AGENT(percept) returns an action
static: KB, a knowledge base
    l, a counter, initially 0, indicating time
TELL(KB, MAKE-PERCEPT-SENTENCE(percept, l))
action ← ASK(KB, MAKE-ACTION-QUERY(l))
TELL(KB, MAKE-ACTION-SENTENCE(action, l))
l ← l + 1
return action
```

### Logic and AI

#### Wumpus:

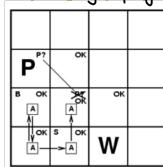
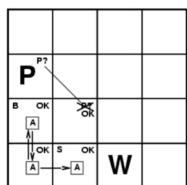
Performance measure: gold +1000; death -1000

→ 1 per step  
→ 10 for using arrow

Sensors: Stench, Breeze, Glitter, Bump, Scream

Actuators: Left turn, Right turn, Forward, Grab, Release, Shoot

**Reasoning turns observation into knowledge**



#### Wumpus in propositional logic:

Let  $P_{i,j}$  be true if there is a pit in  $[i, j]$ .

$\neg P_{1,1}$

Let  $B_{i,j}$  be true if there is a breeze in  $[i, j]$ .

$\neg B_{2,1}$

$B_{2,1}$

$\vdash P_{1,1} ; \vdash \neg B_{2,1}$

$\vdash B_{2,1} \Leftrightarrow (P_{1,2} \vee P_{2,2})$

$\vdash B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$

$\vdash B_{2,1} \Leftrightarrow (W_{1,2} \vee W_{2,2} \vee W_{3,1})$

## Consistency in Knowledge bases

Example: careless while specifying knowledge base:

$$\begin{aligned} \text{PetOfRoommateIsBird} &\Rightarrow \text{PetOfRoommateCanFly} \\ \text{PetOfRoommateIsGryphon} &\Rightarrow \text{PetOfRoommateIsBird} \\ \therefore &\Rightarrow \text{NOT}(\text{PetOfRoommateCanFly}) \end{aligned}$$

Problem: entails PetOfRoommateCanFly and NOT(PetOfRoommateCanFly) → technically, this KB implies anything

→ (logical) KB must be consistent

b "there cannot be contradictions"

## Reasoning Patterns and Modus Ponens:

→ obtain sentences directly from other sentences in KB according to reasoning patterns  
↳ all of logical equivalences from before give us reasoning patterns

Modus Ponens: reasoning pattern → if we have a and  $a \Rightarrow b$ , we can correctly conclude b

## Proof that the sprinklers are on

Knowledge Base	
1) RoommateWet	
2) RoommateWet $\Rightarrow$ (RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers)	
3) RoommateWetBecauseOfSprinklers $\Rightarrow$ SprinklersOn	
4) RoommateWetBecauseOfRain $\Rightarrow$ NOT(RoommateCarryingUmbrella)	
5) UmbrellaGone	
6) UmbrellaGone $\Rightarrow$ (YouCarryingUmbrella OR RoommateCarryingUmbrella)	
7) NOT(YouCarryingUmbrella)	
8) YouCarryingUmbrella OR RoommateCarryingUmbrella (modus ponens on 5 and 6)	
9) NOT(YouCarryingUmbrella) $\Rightarrow$ RoommateCarryingUmbrella (equivalent to 8)	
10) RoommateCarryingUmbrella (modus ponens on 7 and 9)	
11) NOT(NOT(RoommateCarryingUmbrella)) $\Rightarrow$ NOT(RoommateCarryingUmbrella) (equivalent to 10)	
12) NOT(NOT(RoommateCarryingUmbrella)) $\Rightarrow$ NOT(RoommateWetBecauseOfRain) (equivalent to 4 by contraposition)	
13) NOT(RoommateWetBecauseOfRain)	
14) RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers (modus ponens on 1 and 2)	
15) NOT(RoommateWetBecauseOfRain) $\Rightarrow$ RoommateWetBecauseOfSprinklers (equivalent to 14)	
16) RoommateWetBecauseOfSprinklers (modus ponens on 13 and 15)	
17) SprinklersOn (modus ponens on 16 and 3)	

## Rephrasing/Reasoning

## Conjunctive Normal Form (CNF)

→ Any KB can be written as formula in CNF  $\rightarrow (\dots \text{OR} \dots \text{OR} \dots) \text{AND} (\dots \text{OR} \dots \text{OR} \dots) \text{AND} \dots \rightarrow \dots$  can be any literal or its negation → multiple facts ANDed together

e.g. RoommateWet  $\Rightarrow$  (RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers)  $\rightarrow$  (NOT(RoommateWet) OR RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers)

## RoommateWet

- RoommateWet

RoommateWet  $\Rightarrow$  (RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers)

- NOT(RoommateWet) OR RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers

RoommateWetBecauseOfSprinklers  $\Rightarrow$  SprinklersOn

- NOT(RoommateWetBecauseOfSprinklers) OR SprinklersOn

RoommateWetBecauseOfRain  $\Rightarrow$  NOT(RoommateCarryingUmbrella)

- NOT(RoommateWetBecauseOfRain) OR NOT(RoommateCarryingUmbrella)

UmbrellaGone

- UmbrellaGone

UmbrellaGone  $\Rightarrow$  (YouCarryingUmbrella OR RoommateCarryingUmbrella)

- NOT(UmbrellaGone) OR YouCarryingUmbrella OR RoommateCarryingUmbrella

NOT(YouCarryingUmbrella)

- NOT(YouCarryingUmbrella)

## (General) Resolution:

if we have  $l_1 \text{ OR } l_2 \text{ OR } \dots \text{ OR } l_k$

and

$m_1 \text{ OR } m_2 \text{ OR } \dots \text{ OR } m_n$  where for some  $i, j, l_i = \text{NOT}(m_j) \rightarrow$  we can conclude

$l_1 \text{ OR } l_2 \text{ OR } \dots \text{ OR } l_{i-1} \text{ OR } l_{i+1} \text{ OR } \dots \text{ OR } l_k \text{ OR } m_1 \text{ OR } m_2 \text{ OR } \dots \text{ OR } m_{j-1} \text{ OR } m_{j+1} \text{ OR } \dots \text{ OR } m_n$

→ Remove reappearing literals

Knowledge Base	
1) RoommateWet	
2) NOT(RoommateWet) OR RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers	
3) NOT(RoommateWetBecauseOfSprinklers) OR SprinklersOn	
4) NOT(RoommateWetBecauseOfRain) OR NOT(RoommateCarryingUmbrella)	
5) UmbrellaGone	
6) NOT(UmbrellaGone) OR YouCarryingUmbrella OR RoommateCarryingUmbrella	
7) NOT(YouCarryingUmbrella)	
8) NOT(RoommateWet) OR RoommateWetBecauseOfRain OR SprinklersOn (2,3)	Resolution
9) NOT(YouCarryingUmbrella) OR NOT(RoommateWet) OR SprinklersOn (4,8)	
10) NOT(UmbrellaGone) OR YouCarryingUmbrella OR NOT(RoommateWet) OR SprinklersOn (6,9)	
11) YouCarryingUmbrella OR NOT(RoommateWet) OR SprinklersOn (5,10)	
12) NOT(RoommateWet) OR SprinklersOn (7,11)	
13) SprinklersOn (1,12)	

## Horn Clauses: (special case)

→ implications with only positive literals  $x_1 \text{ AND } x_2 \text{ AND } x_4 \rightarrow x_3 \text{ AND } x_5$

True  $\rightarrow x_4$

→ to tell whether  $x_j$  is entailed → follow implications (HP) as far as possible → reach  $x_j$

→ if it can be reached  $x_j$  entailed

→ implement more efficiently by counting reached literals

$(a \Rightarrow b) \equiv (\text{NOT}(b) \Rightarrow \text{NOT}(a)) \rightarrow$  contra position

$(a \Rightarrow b) \equiv (\text{NOT}(a) \text{ OR } b) \rightarrow$  implication elimination

$\text{NOT}(a \text{ AND } b) \equiv (\text{NOT}(a) \text{ OR } \text{NOT}(b)) \text{ OR } (\text{NOT}(a) \text{ AND } \text{NOT}(b)) \rightarrow$  De Morgan

$(a \text{ AND } b \text{ OR } c) \equiv (a \text{ AND } b) \text{ OR } (a \text{ AND } c) \rightarrow$  distributivity

$(a \text{ OR } (b \text{ AND } c)) \equiv ((a \text{ OR } b) \text{ AND } (a \text{ OR } c)) \rightarrow$  associativity

Inference / Entailment: - knowledge base of things we know are true

→ RoommateWetBecauseOfSprinklers  $\Rightarrow$  SprinklersOn

↳ SprinklersOn is entailed by knowledge base if, for every setting of propositional variables for which KB is true (model of knowledge base), SprinklersOn also true

Algorithm to find out if sentence is entailed by KB  
→ Go through possible settings of propositional variables

→ if KB true and a false → return false → else return true  
Problem: not very efficient → Number of Settings:  $2^{\# \text{propositional variables}}$

RWBOS	SprinklersOn	Knowledge base
false	false	false
false	true	false
true	false	false
true	true	true

## Back to the Moon

- 1) PetOfRoommateIsABird  $\Rightarrow$  PetOfRoommateCanFly
- 2) PetOfRoommateIsAPenguin  $\Rightarrow$  PetOfRoommateIsABird
- 3) PetOfRoommateIsAPenguin  $\Rightarrow$  NOT(PetOfRoommateCanFly)
- 4) PetOfRoommateIsAPenguin
- 5) PetOfRoommateIsABird (modus ponens on 4 and 2)
- 6) PetOfRoommateCanFly (modus ponens on 5 and 1)
- 7) NOT(PetOfRoommateCanFly) (modus ponens on 4 and 3)
- 8) NOT(PetOfRoommateCanFly)  $\Rightarrow$  FALSE (equivalent to 6)
- 9) FALSE (modus ponens on 7 and 8)
- 10) FALSE  $\Rightarrow$  TheMoonIsMadeOfCheese (tautology, i.e., always true)
- 11) TheMoonIsMadeOfCheese (modus ponens on 9 and 10)

## Unit resolution

→ implementing general reasoning patterns for normal form

→  $l_1 \text{ OR } l_2 \text{ OR } \dots \text{ OR } l_k$  and  $\text{NOT}(l_i) \rightarrow$  conclude:  $l_1 \text{ OR } l_2 \text{ OR } \dots \text{ OR } l_{i-1} \text{ OR } l_{i+1} \text{ OR } \dots \text{ OR } l_k$

↳ modus ponens

Knowledge Base	
1) RoommateWet	
2) NOT(RoommateWet) OR RoommateWetBecauseOfRain OR RoommateWetBecauseOfSprinklers	
3) NOT(RoommateWetBecauseOfSprinklers) OR SprinklersOn	
4) NOT(RoommateWetBecauseOfRain) OR NOT(RoommateCarryingUmbrella)	
5) UmbrellaGone	
6) NOT(UmbrellaGone) OR YouCarryingUmbrella OR RoommateCarryingUmbrella	
7) NOT(YouCarryingUmbrella)	
8) NOT(UmbrellaGone) OR RoommateCarryingUmbrella (6,7)	Unit Resolution
9) RoommateCarryingUmbrella (5,8)	
10) NOT(RoommateWetBecauseOfRain) (4,9)	
11) NOT(RoommateWet) OR RoommateWetBecauseOfSprinklers (2,10)	
12) RoommateWetBecauseOfSprinklers (1,11)	
13) SprinklersOn (3,12)	

Unit Resolution not enough  $\therefore \rightarrow P \text{ OR } Q, \text{ NOT}(P) \text{ OR } Q \rightarrow$  can we conclude Q?

## Use Resolution for Systematic Inference:

Satisfiable: exists a model that makes modified KB true i.e. mod. KB consistent

Strategy: see if sentence entailed, add NOT(a) to KB and see if it becomes inconsistent

→ CNF for mod. KB satisfiable only if sentence a not entailed

## Resolution Algorithm:

Given: ?formula in CNF

1. Find two clauses with complementary literals

2. Apply resolution

3. Add resulting clause (if not already there)

4. Test: if it results in empty clause, formula not satisfiable

Example: KB: 1) RoommateWetBecauseOfSprinklers

2) NOT(RoommateWetBecauseOfSprinklers) OR SprinklersOn

→ add: 3) NOT(SprinklersOn)

From 2 and 3 get: 4) NOT(RoommateWetBecauseOfSprinklers)

From 4 and 1 get: empty clause

- Limitations of Propositional Logic:
- > some statements hard to model e.g. any (count)
  - > No notion of objects
  - > No notion of relations among objects
  - > Roommate Carrying Umbrella Instructive to us Suggesting:  
objects: Roommate, Umbrella  
relationship between objects: carrying
  - Formally none of this meaning is there → replaceable by any other placeholder names

## 7 Logic and AI 2: First-Order Logic

Elements of First Order Logic (FOL): Objects, Relations, Functions, Equality

functions: → used to encode integers and also data structures

- > could have separate name for Roommate (Person) e.g. Roommate(Person) = Person → not necessary
- > can be applied to any object

video on slide 8

Universal and Existential Quantifiers: → talk about many objects at once

- define variables → refer to one or multiple
- For all:  $\forall$  e.g. all lions are cats  $\forall x: \text{Lion}(x) \Rightarrow \text{Cat}(x)$
- There exists:  $\exists$  e.g. there is a cat that is not a lion  $\exists x: \text{Cat}(x) \Rightarrow \text{NOT}(\text{Lion}(x))$

"John has Jane's umbrella"

- Has(John, Umbrella(Jane))
- "John has an umbrella"
- there exists y: (Has(John, y) AND IsUmbrella(y))
- "Anything that has an umbrella is not wet"
- for all x: ((there exists y: (Has(x, y) AND IsUmbrella(y))) => NOT(IsWet(x)))
- "Any person who has an umbrella is not wet"
- for all x: (IsPerson(x) => ((there exists y: (Has(x, y) AND IsUmbrella(y))) => NOT(IsWet(x))))

"John has at least two umbrellas"

- there exists x: (there exists y: (Has(John, x) AND IsUmbrella(x) AND Has(John, y) AND IsUmbrella(y) AND NOT(x=y)))
- "John has at most two umbrellas"
- for all x, y, z: ((Has(John, x) AND IsUmbrella(x) AND Has(John, y) AND IsUmbrella(y) AND Has(John, z) AND IsUmbrella(z)) => (x=y OR x=z OR y=z))

For all x: a equivalent to NOT (there exists x: NOT(a))

Resolution for FOL:

for all x: (NOT(Knows(John, x)) OR IsMean(x) OR Loves(John, x))

for all y: (Loves(Jane, y) OR Knows(Jane, y))

↳ substitution {x/Jane, y/John}

IsMean(Jane) OR Loves(John, Jane) OR Loves(Jane, John) → Complete (i.e. if not satisfiable, will find proof of this), if we can remove literals that are duplicates after unification → put everything in canonical form first

First-order CNF (canonical form used for resolution):

1. Convert to Negative Normal Form i.e. negation only before predicate symbols

2. Rename variable names used twice within scopes of different quantifiers to names not already used

3. Skolemize Statement: 1. Move quantifiers out so we get "for all x<sub>1</sub>, x<sub>2</sub>... exists y<sub>1</sub>, y<sub>2</sub>..." formula → Quantifiers only appear in initial prefix but never inside NOT, AND, OR  
2. Replace existentially quantified variables by skolem functions  
3. Discard all universal quantifiers

4. Convert into clause set

Example	Steps to CNF (see previous slide)
1. Anything anyone eats and is not killed, is food.	1. Text to FOL
2. For all x: For all y: eats(x,y) AND NOT(killed(x)) => food(y)	2. Eliminate all implications
3. For all x: For all y: NOT(eats(x,y) AND NOT(killed(x))) => food(y)	3. Move Negations inwards
4. For all x: For all y: NOT(eats(x,y)) OR killed(x) OR food(y)	4. Rename variables
5. NOT(eats(x,y)) OR killed(x) OR food(y)	5. Eliminate existential quantifiers (Skolemization)
	6. Drop universal quantifiers
	7. Convert into clauses

Variables start with A, B, ... Z or underscore

Constant "atoms" start with a, b, ... z or in single quotes: curly, 'CS325'

Booleans constants: integers, real numbers, empty list

compatible (Person 1, Person 2) :- eats(Person 1, Food1), eats(Person 2, Food2).

if ? :- eats(Person 1, Food1), eats(Person 2, Food2).

→ Person 1 & 2 compatible if Food exists they both eat

→ to satisfy head of rule → satisfy body

→ compatible (Josie, x) → answer: x = Sam → does not report Food or Person 1 or 2 → local variables

Inference in FOL:

→ Deciding whether sentence is entailed is semidecidable: → algorithms that will eventually produce proof of any entailed sentence

↳ not decidable → cannot always conclude that sentence is not entailed

Program in Logic (Prolog)

Object	Person	Food
Program:	Sam	Curry
eats	Josie	Curry
↳ eats(Sam, Curry)		
eats(Sam, Curry)		
eats(Josie, Curry)		

→ acts like subroutine → type at Prolog prompt eats(Person 1, Food 1) → % constraint over two variables

↳ conjoin two constraints with comma: eats(Person 1, Food 1), eats(Person 2, Food 2)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Sam, Food = Curry → (Person 1 = Sam, Person 2 = Sam, Food = Curry)

↳ solution: Person 1 = Sam, Person 2 = Josie, Food = Curry → (Person 1 = Sam, Person 2 = Josie, Food = Curry)

↳ solution: Person 1 =

- Prolog Solver:** depth - first search using unification and backtracking  
 Query  $\rightarrow$  eats(sam,x)  $\rightarrow$  Iterates in order through program's "eats" clauses  
 $\hookrightarrow$  1st match: eats (sam,dal)  $\rightarrow$  returns  $x=dal$   
 $\hookrightarrow$  "!"  $\rightarrow$  backtracks and continues  $\rightarrow$  next eats (sam,cumy)  $\rightarrow$  returns  $x=cumy$
- $\hookrightarrow$  eats(...,...) facts can be regarded as rows in database  
 ↳ standard relational database operations: eats(x,dal) % select.
- Prolog's backtracking can be inefficient
- video Prolog and Horn Clauses  $\rightarrow$  slide 21
- Applications:** (FOL)
- serious novel mathematical results prove
  - verification of soft- and hardware  $\rightarrow$  prove outputs satisfy required properties for all inputs
  - synthesis of hard- and software  $\rightarrow$  try to prove existence of program satisfying such and such properties  $\rightarrow$  in a constructive way
- Forward and Backward Chaining:
- Forward Chaining: bottom-up approach  $\rightarrow$  starting initial state and reaches goal state, Process of making conclusion based on known facts or data, data-driven we reach goal using available data
- Backward Chaining: top-down approach, based on modus ponens inference rule, goal broken into sub-goal(s) to prove facts true, goal-driven approach  $\rightarrow$  list of goals decides which rules are selected and used, depth-first search strategy for proof
- Something we cannot do in FOL:
- not allowed to reason in general about relations and functions
  - If John is Jack's roommate, then any property of John is also a property of Jack's roommate"  $\rightarrow$  higher order logic  $\rightarrow$  (John = Roommate (Jack))  $\Rightarrow$  for all p: (p(John)  $\Rightarrow$  p(Roommate (Jack)))
  - "If a property is inherited by children, then for any thing, if that property is true of it, it must also be true for any child of it"  $\rightarrow$  for all p: (IsInherentlyChildren(p)  $\Rightarrow$  (for all x,y: ((IsChildOf(x,y) AND p(y))  $\Rightarrow$  p(x))))
- AI Debate 2: link on slide 40
- Symbols, Neurons and Neuro-symbolic:
- Symbol: "A physical symbol system has the necessary and sufficient means for general intelligent action"  $\sim$  Allen Newell & Herbert Simon
- Neurons: "Symbols are luminiferous either" of AI"  $\sim$  Geoff Hinton
- Focus on computational and algorithmic level not just implementational
- Computational: why do things work the way they do? ; what is goal of computation? ; what are unifying principles?
- Algorithmic: What representations can implement such computations? ; how does the choice of representations determine the algorithm?
- Implementational: How can such a system be built in hardware? How can neurons carry out the computations?

Gödel's incompleteness theorem:  $\rightarrow$  video link on slide 32, 84

- FOL cannot model basic arithmetic
- for any consistent system of axioms that can capture basic arithmetic there exist true sentences not provable by axioms
- shows there is no chance to prove every truth in "universe"

Challenging exercise:  $\rightarrow$  another exercise on slide 36

Suppose:  $\rightarrow$  3 objects in world  $\rightarrow$  if x spouse of y, then y spouse of x  
 prove: Something is its own spouse

there exist x,y,z: ( $\neg$ (x=y) AND  $\neg$ (x=z) AND  $\neg$ (y=z))  
 for all w, x,y,z: (w=x OR w=y OR w=z OR x=y OR x=z OR y=z)  
 for all x,y: ((spouse(x)=y)  $\Rightarrow$  (spouse(y)=x))  
 for all x,y: ((spouse(x)=y)  $\Rightarrow$   $\neg$ (x=y))  $\rightarrow$  contradiction



## 8 Uncertainty $\rightarrow$ agents have to deal with uncertain situations / uncertainty

$\hookrightarrow$  qualification problem  $\rightarrow$  impossibility to model all things that can go wrong

Probabilities:  $\rightarrow$  way of handling uncertainty

$\hookrightarrow$  may summarize effects that are due to laziness (don't want to list everything that can go wrong), theoretical ignorance (some things cannot be known), practical ignorance (things might not be known about situation)

$\rightarrow$  can be related to (subjective) beliefs:  $\rightarrow$  probability p means I believe statement will be true in p...100% of cases  $\rightarrow$  e.g. traffic in 10% of cases  $\rightarrow$  not traffic by degree of 10%?

probability theory is about the degree of belief not degree of truth

$\hookrightarrow$  probability of propositions change with new evidence (e.g.  $P(A_{15})$  gets me there in time | no reported accidents) = 0.05 but  $P(A_{15})$  will get me there | no reported accidents, 50m) = 0.15

$\hookrightarrow$  Consider probability that sun is still there tomorrow  $\rightarrow$  difficult to observe by experiment

$\hookrightarrow$  Does patient have illness  $\rightarrow$  compare to other patients, but don't take too many information into consideration or else very hard to compare

Basics:

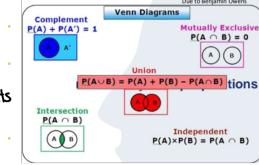
state or sample space  $\rightarrow$  set of all samples (e.g. die  $\rightarrow$  1,2,3,4,5,6)  $\Omega$   $w \in \Omega$

probability space / model  $\rightarrow$  sample space with assignment of probabilities per possible sample (e.g.  $P(1)=P(2)=P(3)=\dots=P(6)=\frac{1}{6}$ )  $P(w)$  for every  $w \in \Omega$ ;  $\sum_w P(w) = 1$

event  $\rightarrow$  subset of sample space (e.g.  $P(\text{roll greater } 4) = P(5) + P(6) = \frac{1}{6} + \frac{1}{6} = \frac{1}{3}$ )  $P(A) = \sum_{w \in A} P(w)$

Kolmogorov's Axioms of Probability:

1. All probabilities between 0 & 1  $\rightarrow 0 \leq P(a) \leq 1$
2. Necessarily true propositions have probability 1, ..., have probability 0  $\rightarrow P(\text{false}) = 0, P(\text{true}) = 1$
3. Probability of disjunction is  $P(a \vee b) = P(a) + P(b) - P(a \wedge b)$
4. Axioms restrict set of probabilistic beliefs a agent can (reasonably) hold  $\rightarrow$  similar to logical constraints



## Example: why you shouldn't violate the axioms of probability:

Example:  
→ suppose Agent 1 believes the following:  $P(a) = 0.4$ ,  $P(b) = 0.3$ ,  $P(a \vee b) = 0.8$

Agent 2 can now select a set of events and bet on them according to these probabilities so that she cannot loose (e.g. Agent 2 offers e.g. \$6 against Agent 1's \$4 for proposal a)

Agent 1		Agent 2		Outcome for Agent 1			
proposal	belief	state	bet	a	b	a $\vee$ b	-a $\wedge$ -b
Axioms of Probability				4	6	6	4
axiom 1: $P(a) + P(b) = 1$				7	3	7	3
axiom 2: $P(a \vee b) = P(a) + P(b)$				2	2	2	2
Agent 2 now designs a game, choosing to bet \$4 on a, \$3 on b and \$2 on not(a $\wedge$ b). E.g., if a holds, Agent 2 would have to pay \$6, i.e. -6 \$2 on not(a $\wedge$ b). So summing up, Agent 1 always loses money				-11	-1	-1	-1

Random variable  $\rightarrow$  function from atomic events to some range of values

e.g. Roulette  $\rightarrow$  Atomic event: 0-36

Random variables with outcomes true or false: Rouge/Mair, Pair/Impair, Passee/Fausse  
Transversale, Carré, Cheval

$\hookrightarrow$  range (36) = true

probability function  $P$  over atomic events induces a probability distribution over all random variables  $X$   
 $\hookrightarrow P(X=x_i) = \sum_{\omega: X(\omega)=x_i} P(\omega) = P(x_i)$

Joint Distributions:  $\rightarrow$  gives probability of combined events

(e.g. probability that  $X=x$  and  $Y=y \rightarrow$  true:  $P(X,y) = P(X=x \wedge Y=y)$ )

$\rightarrow$  Reduce complexity by making use of independencies

$X \& Y$  independent from another if:  $P(X|Y) = P(X)$ ,  $P(Y|X) = P(Y)$ ,  $P(X,y) = P(x)P(y)$ .

$\hookrightarrow$  not affected by the other variable  $\rightarrow$  reduces amount of possible values

$\rightarrow$  Absolute independent variables are rare i.e. no variables in cancer research are independent

Conditional Probabilities:  $\rightarrow$  probability of  $X=x$  under assumption  $Y=y$  is true

$P(X|Y) = \frac{P(x \wedge y)}{P(y)}$   $\rightarrow$  alternative formulation Product rule  $\rightarrow P(X,y) = P(X|y)P(y) = P(y|x)P(x)$

e.g. assume  $P(Y=few) = P(Y=few) = 0.15$  calculate  $P(X=maligne|Y=few) = \frac{P(maligne,few)}{P(few)} = \frac{0.005}{0.15} = 0.04$   
assume  $P(X=maligne)$  calculate  $P(maligne) = 0.019$  calculate  $P(Y=few|X=maligne) = \frac{P(maligne,few)}{P(maligne)} = \frac{0.005}{0.019} = 0.316$

Bayes Rule: Probability of evidence, given belief is true. Likelihood  $\rightarrow$

$$P(X|Y) = \frac{P(Y|X)P(X)}{P(Y)} \quad \begin{array}{l} \text{Probability of hypothesis} \\ \text{before considering evidence.} \\ \text{Prior} \end{array}$$

$\begin{array}{l} \text{Probability of evidence } Y \\ \text{under any circumstance} \\ \text{Marginalization} \end{array}$

Propositions, or towards uncertain knowledge:

$\rightarrow$  Often in AI applications sample points are defined by values of a set of random variables

$\hookrightarrow$  i.e. Sample Space is Cartesian product of ranges of variables

$\rightarrow$  With Boolean variables, sample points = propositional logic model (e.g. A=true, B=false, or A  $\wedge$  B)

$\rightarrow$  Proposition = disjunction of atomic events in which it is true (e.g.  $(a \vee b) = (\neg a \wedge b) \vee (a \wedge \neg b) \vee (a \wedge b) \rightarrow P(a \vee b) = P(\neg a \wedge b) + P(a \wedge \neg b) + P(a \wedge b)$ )

Syntax:

$\rightarrow$  Propositional or Boolean random can be true or false (e.g. hasFluBacteria (proposition))

$\rightarrow$  Discrete random variables (finite or infinite) (e.g. Weather one of sunny, rain, cloudy, snow)

$\rightarrow$  Weather=rain  $\rightarrow$  proposition; values must be exhaustive and mutually exclusive

$\rightarrow$  Continuous random variables (bound or unbound) (e.g. Temp is unbound variable)

$\rightarrow$  Temp > 25.5, Temp > 23 are propositions

Cancer

Marginalization:

$\forall$  any set of  $X$  and  $Y \rightarrow$  compute probability

$\hookrightarrow P(Y) = \sum_{i=1}^n P(x_i, Y)$

$\rightarrow$  resulting distribution  $\rightarrow$  marginal distribution

$\hookrightarrow$  its probabilities are marginal probabilities

no	benign	maligne
0.768	0.024	0.008
0.132	0.012	0.006
0.035	0.010	0.005

$$P(Y=few) = P(\text{no,few}) + P(\text{benigne,few}) + P(\text{maligne,few})$$

$$\begin{aligned} \text{Chain rule derived by successive application of Product rule: } & P(X_1, X_2, \dots, X_n) = \\ & P(X_1, \dots, X_{n-1}) P(X_n | X_1, \dots, X_{n-1}) = P(X_1, \dots, X_{n-2}) P(X_{n-1} | X_1, \dots, X_{n-2}) P(X_n | X_1, \dots, X_{n-2}) \\ & = \prod_{i=1}^n P(x_i | X_1, \dots, X_{i-1}) \end{aligned}$$

e.g. AIDS-test: Events: AIDS = infected or not

Positive = positive test result

Probabilities:  $P(\text{positive}, \text{AIDS}) = 0.99$

$P(\text{negative}, \text{AIDS}) = 0.01$

$P(\text{positive}, \text{notAIDS}) = 0.005$

$P(\text{negative}, \text{notAIDS}) = 0.995$

$$\begin{aligned} \text{assume } & P(\text{AIDS}) = 0.0001 \text{ (Prior)} + \text{positive test} \\ & P(\text{AIDS}) = \frac{P(p|\text{A})P(\text{A})}{P(p)} = \frac{P(p|\text{A})P(\text{A})}{P(p|\text{A})P(\text{A}) + P(p|\text{notA})P(\text{notA})} \\ & = \frac{0.99 \cdot 0.0001}{0.99 \cdot 0.0001 + 0.005 \cdot 0.9999} = 0.0194 \end{aligned}$$

Nobel Prize Economics (2002, Amos Tversky & Daniel Kahneman): Uncovered number of biases seeming to characterize human reasoning & decision-making, providing significant challenge to economic models assuming people simply apply statistical decision theory.

Uncertainty in AI: Recall joint distribution is enumerating everything  $\rightarrow$  Worst-case run time:  $O(2^n)$  ( $n = \# \text{ of RVs}$ ), Space is  $O(2^n) \rightarrow$  size of joint distribution's table

$\hookrightarrow$  Main idea: make use of independencies to compress representation

## 9 Bayesian networks

$\rightarrow$  simple, graphical notation for conditional independence assertions, hence for compact specifications of full joint distributions  
 $\rightarrow$  directed acyclic graph with  $\rightarrow$  Nodes (one for each variable) and Edges (directed edge from node  $N_i$  to node  $N_j$  to indicate  $N_i$  has direct influence upon  $N_j$ )

Naïve Bayes model:  $\rightarrow$  assumes all effects are independent given cause

$$P(\text{hypothesis}, \text{evidence}_1, \text{evidence}_2, \dots, \text{evidence}_n) = P(\text{hypothesis}) \prod_i P(\text{evidence}_i | \text{hypothesis})$$

total # of parameters is linear in  $n$ .  
Graphical encoding of conditional distributions.

Set of random variables  $\{X_1, \dots, X_n\}$

Directed acyclic graph (DAG)

To each RV  $X_i$  associate conditional probability distribution:  $P(X_i | \text{Pa}(X_i))$

joint distribution  $P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i | \text{Pa}(X_i))$

BN semantics



$$\begin{aligned} P(G, A) &= P(G|A) \cdot P(A) \\ P(A|G) &= P(A) \\ P(G|A) &= P(G) \end{aligned}$$

Absolute independencies are rare  
 $\hookrightarrow$  less entries and consequently lower complexity  
 $P(C|S, G, A) = P(C|S)$

e.g. Smoking  $\rightarrow$  Cancer

$$S \in \{\text{no, few, many}\} \quad C \in \{\text{no, benign, maligne}\}$$

	n	f	m
P(S=n)	0.80		
P(S=f)	0.15		
P(S=m)	0.05		
P(C=n)	0.96	0.88	
P(C=f)	0.03	0.08	0.25
P(C=m)	0.01	0.04	0.15

Local Markov Assumption

Each RV  $X_i$  independent of its "non-descendants" given its parents ( $X_i \perp \text{non-descendants} | \text{Pa}(X_i)$ )

Interference:

$$\text{Query: } P(X|e)$$

Definition of conditional probability  $P(X|e) = \frac{P(X,e)}{P(e)}$

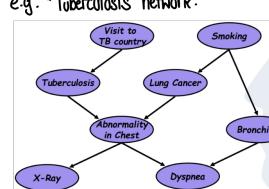
Up to normalization  $P(X|e) \propto P(X,e)$

$\hookrightarrow$  rewrites to  $P(Y) = \sum_{X_i \in e} P(X_i; \text{Pa}(X_i))$

BN semantics

$$\Sigma_a (P_a \cdot P_e) = (\Sigma_a P_a) \cdot P_e \quad \text{Marginalization}$$

$\hookrightarrow$  if A not in  $P_e$



**Variable Elimination:**

- compute  $P(d)$
- need to eliminate:  $v, s, x, t, l, a, b$

$$P(v)P(s)P(t|v)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

Eliminate:  $v$

$$\text{Compute: } f_v(t) = \sum_s P(s)P(t|v)$$

$$\Rightarrow f_v(t)P(s)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

Note:  $f_v(t) = P(t)$   
In general, result of elimination is not necessarily a probability term

$$P(v)P(s)P(t|v)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)P(s)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

Eliminate:  $s$

$$\text{Compute: } f_s(b,l) = \sum_s P(s)P(b|s)P(l|s)$$

$$\Rightarrow f_s(b,l)P(a|t,l)P(x|a)P(d|a,b)$$

Summing on  $s$  results in a factor with two arguments  $f_s(b,l)$   
In general, result of elimination may be a function of several variables

$$P(v)P(s)P(t|v)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)P(s)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)f_s(b,l)P(a|t,l)P(x|a)P(d|a,b)$$

Eliminate:  $x$

$$\text{Compute: } f_x(a) = \sum_x P(x|a)$$

$$\Rightarrow f_x(a)f_s(b,l)f_a(a)P(a|t,l)P(d|a,b)$$

Note:  $f_x(a) = 1$  for all values of  $a$  !!

$$P(v)P(s)P(t|v)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)P(s)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)f_s(b,l)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)f_s(b,l)f_a(a)P(a|t,l)P(d|a,b)$$

Eliminate:  $t$

$$\text{Compute: } f_t(a,l) = \sum_t f_v(t)P(a|t,l)$$

$$\Rightarrow f_s(b,l)f_a(a)f_t(a,l)P(d|a,b)$$

$$P(v)P(s)P(t|v)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)P(s)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)f_s(b,l)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)f_s(b,l)f_a(a)P(a|t,l)P(d|a,b)$$

$$\Rightarrow f_s(b,l)f_a(a)f_t(a,l)P(d|a,b)$$

Eliminate:  $/$

$$\text{Compute: } f_{/}(a,b) = \sum_l f_t(a,l)f_a(a)$$

$$\Rightarrow f_s(b,l)f_a(a)P(d|a,b)$$

$$P(v)P(s)P(t|v)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)P(s)P(l|s)P(b|s)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)f_s(b,l)P(a|t,l)P(x|a)P(d|a,b)$$

$$\Rightarrow f_v(t)f_s(b,l)f_a(a)P(a|t,l)P(d|a,b)$$

$$\Rightarrow f_s(b,l)f_a(a)f_t(a,l)P(d|a,b)$$

$$\Rightarrow f_s(b,l)f_a(a)P(d|a,b) \rightarrow f_b(b,d) \rightarrow f_d(d)$$

Eliminate:  $a, b$

$$\text{Compute: } f_a(b,d) = \sum_a f_s(b,a)f_a(a)P(d|a,b) \quad f_b(d) = \sum_b f_s(b,d)$$

**Variable elimination:**

- Given: BN and query  $P(X|e) / P(X,e)$
- Instantiate evidence  $e$
- Choose elimination order over variables, e.g.  $X_1, \dots, X_n$
- Initial factors  $\{f_1, \dots, f_n\}$ :  $f_i = P(X_i | P_{X_i})$  (CPT for  $X_i$ )
- For  $i=1$  to  $n$ , if  $X_i \notin \{X_j\}$

  - Collect factors  $f_1, \dots, f_k$  that include  $X_i$
  - Generate new factor by eliminating  $X_i$  from these factors  

$$g = \sum_{X_i} \prod_{j=1}^k f_j$$
  - Variable  $X_i$  has been eliminated! Add  $g$  to set of factors

Normalize (everything sums to 1)  $P(X,e)$  to obtain  $P(X|e)$

**Theorem:** Inference (exact/approximate) in BN is NP-hard (#P; via reduction to 3-SAT)

$$(\bar{x}_1 \vee x_2 \vee x_3) \wedge (\bar{x}_2 \vee x_3 \vee x_4) \dots$$

$$C_1 \quad C_2 \quad \dots \quad C_n \quad Y$$

**Approximate Inference:** → by stochastic sampling (from a BN)

- Draw  $N$  samples from sampling distribution  $S$
- Compute approximate posterior probability  $\hat{P}$
- Show this converges to the true probability  $P$

**Outline:**

- Sampling from empty network
- Rejection sampling: reject samples disagreeing with evidence
- Likelihood weighting: use evidence to weight samples
- Markov Chain Monte Carlo (MCMC): Sample from a stochastic process whose stationary distribution is true posterior

**Draw Sample:**

Given: Random variable  $X, \mathcal{P}(X) = \{0,1\}, P(x) = \{0.3, 0.7\}$   
Sample  $X \leftarrow \mathcal{P}(X)$ : Draw random number  $r \in [0,1]$ , If  $r < 0.3$  set  $X=0$ , else  $X=1$  → can generalize of any domain size

Sampling from an "Empty Network": - Generating Samples from network having no evidence associated with it

- sample value for each variable in topological order
- using the specified conditional probabilities

**function** PRIOR-SAMPLE( $bn$ ) **returns** an event sampled from  $bn$

**inputs:**  $bn$ , a belief network specifying joint distribution  $P(X_1, \dots, X_n)$

$x \leftarrow$  an event with  $n$  elements

for  $i = 1$  to  $n$  do

$x_i \leftarrow$  a random sample from  $P(X_i | \text{parents}(X_i))$   
given the values of  $\text{Parents}(X_i)$  in  $x$

**return**  $x$

**e.g.**

**Probability Estimation:**

- Sample many points using algorithm above
- count how often each possible combination  $x_1, x_2, \dots, x_n$  appears
- estimate probability by observed percentages  $\hat{P}_{(x_1 \dots x_n)} = N_{(x_1 \dots x_n)} / N_{\text{total}}$
- ↳ converges towards joint probability function!

• What to do when we find a problem that looks hard.  
I couldn't find a polynomial-time algorithm; I guess it's too hard.

• Sometimes we can prove a strong lower bound... (but not usually)  
I couldn't find a polynomial-time algorithm, because no such algorithm exists!

• NP-completeness let's us show collectively that a problem is hard.  
I couldn't find a polynomial-time algorithm, but neither could all these other smart people.

NP-Complete

NP

NP-hard

NP

poly-time

EXPSPACE

EXPTIME

PSpace

NP

P

NL

Computers and Intractability: A Guide to the Theory of NP-Completeness

Michael R. Garey, David S. Johnson

0.5/0.5 prior

$P(Y=1) > 0$  iff 3-SAT formula is satisfiable

## Markov Chain Monte Carlo (MCMC) Sampling:

- "State" of network = current assignment to all variables
- Generate next state by sampling one variable given Markov blanket
- Sample each variable in turn, keeping evidence fixed

```
function MCMC-Ask(X; e, N)
    local variables:  $N[X]$ , a vector of counts over  $X$ , initially zero
    Z: the nonevidence variables in  $N$ 
    x: the current state of the network, initially copied from e
    initialize x with random values for the variables in Y
    for j = 1 to N do
        for each  $Z_i$  in Z do
            sample the value of  $Z_i$  in x from  $P(Z_i | \text{Ind}(Z_i))$ 
            update the value of  $M(Z_i)$  in x
             $N[x] = N[x] + 1$  where  $x$  is the value of  $Z_i$  in x
    return NORMALIZE( $N[X]$ )
```

▷ Can also choose a variable to sample at random each time

### Ordered Gibbs Sampler:

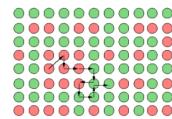
▷ Generate sample  $x^{+M}$  from  $x^+$ : Process all variables in some order

$$\begin{aligned} X_1 &= X_1^{+M} \leftarrow P(x_1 | x_2^+, x_3^+, \dots, x_N^+, e) \\ X_2 &= X_2^{+M} \leftarrow P(x_2 | x_1^{+M}, x_3^+, \dots, x_N^+, e) \\ \vdots & \\ X_N &= X_N^{+M} \leftarrow P(x_N | x_1^{+M}, x_2^{+M}, \dots, x_{N-1}^{+M}, e) \end{aligned}$$

in short, for  $i=1$  to  $N$ :  $X_i = x_i^{+1} \leftarrow \text{sampled from } P(x_i | x_i^+ \setminus x_i; e)$

### Gibbs Sampling: Illustration

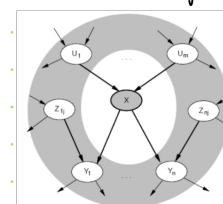
The process of Gibbs sampling can be understood as a random walk in the space of all instantiations with  $Y = u$ .



Reachable in one step: instantiations that differ from current one by value assignment to at most one variable (assume randomized choice of variable  $X_k$ ).

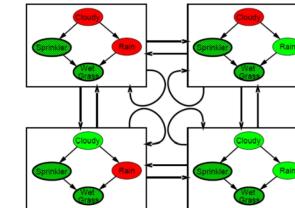
Guaranteed to converge iff chain is:

- irreducible (every state reachable from every other state)
- aperiodic (returns to state i can occur at irregular times)
- ergodic (returns to every state with probability 1)



### Markov Chain:

With  $\text{Sprinkler} = \text{true}$ ,  $\text{WetGrass} = \text{true}$ , there are four states:



Wander about for a while, average what you see

### Get Probability Distribution from Sampling:

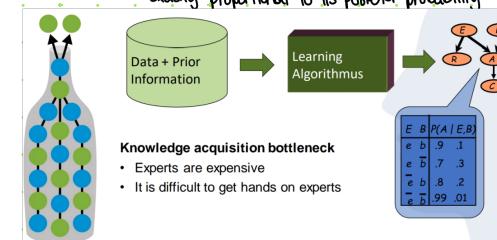
e.g. Task: Estimate  $P(\text{Rain} | \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$

1. Sample Cloudy or Rain given Markov Blanket, repeat n times
2. Count # of times, Rain is true and false in samples

$\rightarrow 1000 \text{ states} \rightarrow 84 \text{ true} \rightarrow 69 \text{ false}$  (Haha :))

$$P(\text{Rain} | \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true}) = \text{Normalize}(84, 69) = (0.84, 0.69)$$

Theorem: Chain approaches stationary distribution:  
long-run fraction of time spent in each state is exactly proportional to its posterior probability



## 10 Machine Ethics

↳ nicht klausurrelevant → lesen / Definitionen anschauen

## 11 Introduction into Machine Learning and Neural Networks

### Learning:

1. essential for dealing with unknown environments (agent not omniscient?)
2. useful as system construction method (expose agent to reality)
3. modifies agent decision mechanism to improve performance  
("insuring" doing same thing over and over again and expecting different results)

Methods → Memorization (Declarative Knowledge)

- ↳ Accumulation of individual facts
- ↳ Limited by time to encode & memory
- Generalization (Imperative Knowledge)
- ↳ Deduce new from old facts
- ↳ Limited by accuracy of deduction process  
(predictive activity, assumes relations between past and future)

Inductive Learning: → simplest form of learning. It learns function from examples.

- f is (unknown) target function we want to learn.  $f(x)$  is target, label or y
- Examples defined as  $(x, f(x)) \rightarrow (x, \text{Rain})$

Problem: find hypothesis h, such that h is f

- ↳ Given training set of examples
- ↳ On all examples

Example:

- Day 1: 10h
- Day 2: 10h →  $f(2) = 10$
- Day 3: 9h
- Day 4: 2h
- ... Day 31: 5h →  $f(32) = 6$

Predicting the future: → construct "rule set" h to agree with f on training set

→ h consistent → agrees with f on all examples

Ockham's Razor → best explanation is simplest explanation fitting the data

Oversizing Avoidance → Maximize combination of consistency and simplicity

### 1 Machine Learning:

→ Program algorithm to automatically learn a program from data or experience.



→ ML often used as synonym for AI → not the same!

- ↳ AI not always imply learning-based system (Search, CSP, logical inference, ...)
- ML focuses on learning-based systems while often extracting knowledge from data

Human learning → (often) data- and knowledge-efficient; complete multitasking, multi-modal system; time-invariant  
ML inspired by human learning → not goal to rebuild human learning  
→ borrow ideas from biological systems (e.g. neural networks); may perform better/worse than humans (takes long)

Applications: Web Search, Computational Biology, Speech Recognition → Machine Translations; Image Recognition, Robotics, Finance and Stock Market, Medical Diagnosis, Information Extraction → Visual Analytics, Traffic Prediction, Software development, ...

Designing learning System: ① Do we need learning approach for problem? (pattern?, analytical solving?, training data?)

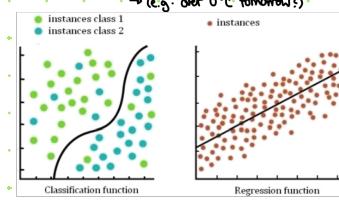
② Type of problem? (representation? Algorithm choosing?) ③ Gather and organize data. (processing important)

④ Fitting/Training model ⑤ Optimization ⑥ Evaluate and iterate back to ② → ML not solution to all problems

## 2 Get to a solution:

### Types of Machine learning:

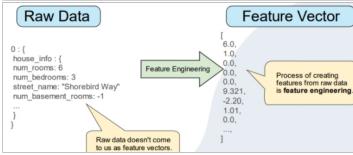
- Supervised Learning (based on labeled datasets, learns to map inputs and outputs based on pairs in dataset used in learning process)
- Unsupervised Learning (unlabeled data → searches for patterns and similarities in data)
- Reinforcement Learning (in RL agent learns by interacting with its environment and getting positive or negative reward)
- Supervised Learning: Classification and Regression Tasks**
- Given dataset:  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n) \rightarrow$  goal: learn function  $h(x)$  to predict  $y$  given  $x$
- Regression task:  $\rightarrow y$  is continuous value, e.g. temperature tomorrow
- Classification task:  $\rightarrow y$  is discrete class label → algorithm tries predicting continuous value describing label probability
- ↳ Propose: over 0°C tomorrow?



Model: representation of what algorithm learned from data, used in training process. Model is output representation of learned "rule set"

Feature Engineering: process of selecting, manipulating and transforming raw data into features used within learning approach

Example	Weather	Location	Date	Temperature	Precipitation
1	Rainy	Darmstadt	11.04.22	12.3	44.2
2	Sunny	Hamburg	11.04.22	19.2	12.6
3	Rainy	Darmstadt	12.04.22	16.7	67.3
4	Cloudy	Heidelberg	11.04.22	17.3	22.2
...	...	...	...	...	...



Know data: distribution?, outliers?, reflect reality?, biased?, ...  
→ using bad data results in bad models → noise per se might not be a problem

### ML Common Approaches:

Regression: Linear Regression, Multiple Linear Regression, Regression trees, Non-linear Regression, Polynomial Regression,

Classification: Random Forest, Decision Trees, Logistic Regression, Naive Bayes, Support Vector Machines, ...

→ no single best model working best for all problems.

### 5 Prediction or Forward Propagation:

Assume input:  $x_1 = 6, x_2 = 4$

Weights:  $w_{0,1}^{(4)} = 1, w_{0,2}^{(4)} = -1, w_1^{(4)} = -1$

$w_2^{(4)} = 0.33, w_{2,1}^{(4)} = 0.5, w_{2,2}^{(4)} = -0.8$

$w_3^{(4)} = 0.1, w_{3,1}^{(4)} = 0.6, w_{3,2}^{(4)} = 2$

Assume bias  $w_{0,0}^{(4)} = 0$

$\Rightarrow z_4 = g(w_{0,0}^{(4)} + \sum_{j=1}^2 x_j \cdot w_{j,0}^{(4)}) = g(0 \cdot 1 + 6 \cdot -1) = g(-6) = g(-1)$

$z_5 = g(z_2 + 2) = g(4)$

$z_6 = g(0.6 + 2 \cdot 4) = g(3)$

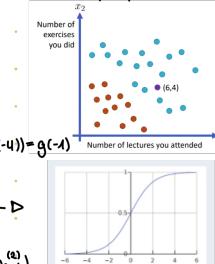
Use sigmoid activation function:  $\sigma(x) = \frac{1}{1 + e^{-x}}$

$\Rightarrow z_4 = 0.26, z_5 = 0.96, z_6 = 0.95$

$\hat{y} = \sigma(w_{0,1}^{(4)} + \sum_{j=1}^2 x_j \cdot w_{j,1}^{(4)}) = \sigma(0 \cdot 1 + 6 \cdot -0.5 + 4 \cdot 0.5) = \sigma(0.5)$

$\hat{y} = \sigma(-1 \cdot 0.26 + (-0.8) \cdot 0.96 + 2 \cdot 0.95) = \sigma(0.856) = 0.7 \rightarrow$  Yes purple should pass

e.g. Will purple pass exam?



### 6 Training a NN or Backpropagation:

Assume instead of 0.7 → 0.14 (fails exam) → How do we measure quality?

Loss function: Quantifying Loss  $J(f(x^{(i)}; w), y^{(i)}) \rightarrow$  cost of incorrect predictions

Empirical Loss  $J(w) = \frac{1}{n} \sum_{i=1}^n J(f(x^{(i)}; w), y^{(i)}) \rightarrow$  Measures total loss over dataset

↳ Objective function, cost function, empirical risk

$$\text{Compute } \frac{\partial J(w)}{\partial w} = \frac{\partial J(w)}{\partial w_1} + \frac{\partial J(w)}{\partial w_2} + \dots + \frac{\partial J(w)}{\partial w_n} \quad \text{Backpropagation}$$

→ How much do weights affect outcome i.e. final loss?  $\frac{\partial J(w)}{\partial w_i}$

Using chain rule describing problem:  $\frac{\partial J(w)}{\partial w_i} = \frac{\partial J(w)}{\partial z_1} \cdot \frac{\partial z_1}{\partial w_i} + \frac{\partial J(w)}{\partial z_2} \cdot \frac{\partial z_2}{\partial w_i} + \dots + \frac{\partial J(w)}{\partial z_n} \cdot \frac{\partial z_n}{\partial w_i}$

↳ repeat for every weight in network using gradients from later layers

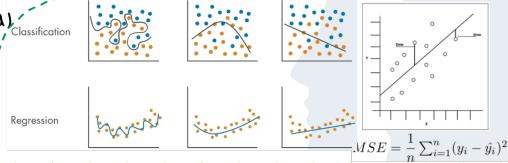
→ Propagate the error back to all nodes, through the network

### 3 Evaluating model: → to measure quality → know goal!

→ Accuracy, Precision, Recall, Mean Squared Error, ...

→ choosing correct metric depends on goal

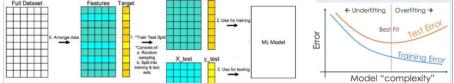
Mean Squared Error



ML Overfitting: model has trained "too well" → memorized dataset while losing ability to generalize i.e. perform on unknown / new data.

Deal with overfitting: Split data, Regularization, Use more data → augment data i.e. adding noise, Select different features, Cross-validation, Ensemble methods...

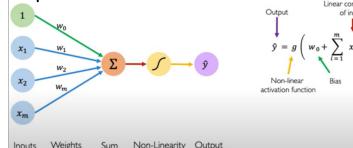
Train / Test Split: → model not only perform on data used in training → Split data into Training & testing



### 4 Artificial Neural Networks:

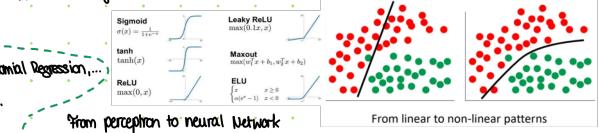
Deep learning: learn underlying features directly from data without specifying them

#### Perception → An Artificial Neuron



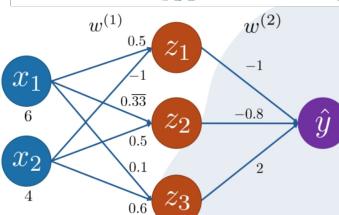
→ Output activation determined by activation function  $g$

Activation functions: → decides if neuron should be active → nowadays mostly non-linear function  
→ adds non-linearity to neural network, without it simple linear regression model, allows use of backpropagation



From perception to neural network

→ Perceptions may have multiple output nodes → output nodes can be combined with other perceptions  
→ Build networks of these nodes i.e. Multilayer Perceptrons (MLP) → information flow unidirectional  
→ Information distributed and processed in parallel  
→ Use size (i.e. # of layers) to model expressiveness of MLP  
→ Following definition of perception we know for each hidden node  $z_{k,i} = \sigma(w_{0,i}^{(k)} + \sum_{j=1}^{k-1} z_{j,i} \cdot w_{j,i}^{(k)})$



#### Training Network:

Overall goal: Minimize loss  $N^* = \arg \min_w \frac{1}{n} \sum_{i=1}^n L(f(x^{(i)}; w), y^{(i)})$

Gradient descent

Algorithm:

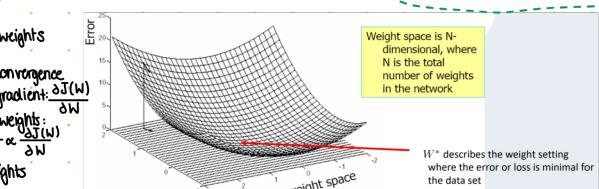
1. Initialize weights randomly

2. Loop until convergence

→ Compute gradient:  $\frac{\partial J(w)}{\partial w}$

→ Update weights:  $w \leftarrow w - \alpha \frac{\partial J(w)}{\partial w}$

3. Return weights



Weight space is N-dimensional, where N is the total number of weights in the network

$w^*$  describes the weight setting where the error or loss is minimal for the data set

**Example:**  $4 \cdot x_1 = 0.5$

Given:  $g(x) = \text{ReLU}(x) = \max(0, x)$ ,  $x = 4$ ,  $y = 1$ ,  $\hat{y} = 2$ ,  $L = -(\hat{y} - y)^2$ ,  $J(W) = L$

Wanted:  $\frac{\partial J(W)}{\partial w_i} = \frac{\partial J(W)}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial w_i}$

**Step 1:**  $\frac{\partial J(W)}{\partial \hat{y}} = 2(\hat{y} - y) = 2(2 - 1) = 2$

**Step 2:**  $\frac{\partial \hat{y}}{\partial w_i} = \text{ReLU}'(w_0 + \sum_{i=1}^n w_i x_i) \cdot x_i = 1 \cdot 4 = 4$

**Step 3:**  $\frac{\partial J(W)}{\partial w_i} = \frac{\partial J(W)}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial w_i} = 2 \cdot 4 = 8$

**Update the weights:**  $W \leftarrow W - \alpha \frac{\partial J(W)}{\partial W}$

**Step 4:**  $W_{\text{new}} = W_{\text{current}} - \alpha \frac{\partial J(W)}{\partial W}$

$W_{\text{new}} = W_{\text{current}} - \alpha \cdot 8$

$W_{\text{new}} = 0.5 - 0.05 \cdot 8$

$W_{\text{new}} = 0.1$

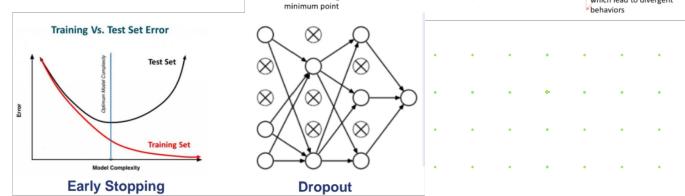
**Alternative Step 4:**  $W_{\text{new}} = 0.5 - 0.5 \cdot 8$

$W_{\text{new}} = -3.5$

→ explanation video link slides 51-53

► Loss function can be difficult to optimize → Hard to find global minimum  
 → Change weights into direction of steepest descent or error function giving step size → learning rate ( $\alpha$ )  
 ↳ difficult finding a good learning rate

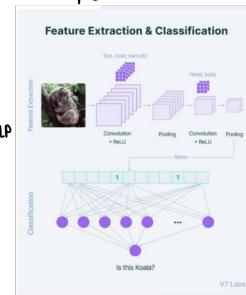
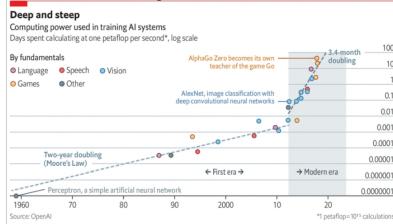
**Regularization:** set of techniques to prevent overfitting in neural networks  
 → improve accuracy of a Deep Learning model when facing completely new data from problem domain



7 What's next:  
 Further Architectures: MLP only first step in field of neural network architectures → How many layers? Complex inputs (images)? # features use? Use images and text as inputs at the same time? How can I do sequences of data e.g. sentences, time-series, ...?

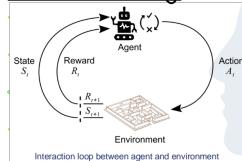
**Convolutional Neural Network (CNN):** uses convolutional layers to extract features from input  
 → Features in first layers refer to edges, borders, shapes, ...  
 → higher levels: pattern detection up to specific objects  
 → uses pooling layers to decrease computational requirements and extract more dominant features  
 → Compared to MLPs: fewer connections, i.e. weights, easier to train, can have a lot of layers, popular in fields like Computer Vision and NLP

## Performance Processing:



## 12 Introduction into Reinforcement Learning and AlphaZero

### Reinforcement Learning:



- RL algorithms attempt to find policy (what action in which state) for maximizing cumulative reward for agent over course of problem
- Mathematically, typically represented by a Markov Decision Process (MDP)
- RL differs from supervised learning in that correct input/output pairs are never presented, nor sub-optimal actions explicitly corrected

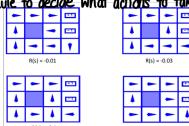
Determine how system's performance success due to contribution of components  
 In RL, i.e. temporal CAP: → Given sequence of actions/states, final sum of time-discounted future rewards, how infer which actions were effective at producing lots of reward and which actions were not effective?  
 → How assign credit for observed Rewards given sequence of actions over time?  
 → Every reinforcement learning algorithm must address this problem

### Example: Fruit Fly

- agent lives in a grid & walls block agent's path
- agent's actions are noise i.e. don't always go as planned
- ↳ 80% action North takes agent North  
 ↳ 10% " West, 10% East (Stochastic movement)
- ↳ if there is wall → no movement
- small "living" reward each step → big rewards come at the end → Goal: maximize sum of rewards
- Deterministic grid world: 100% taking action supposed to take
- Stochastic grid world: because of noise multiple actions could happen

### Solve MDPs: Finding optimal plan

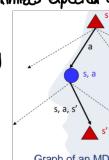
- in deterministic single-agent search problems we want optimal plan, or sequence of actions, from start to goal
- in MDP compute optimal policy  $\pi^*: S \rightarrow A$ : requires an action for each state, optimal policy maximizes expected utility if followed, defines reflex agent
- Policy  $\pi_1$ : Rule to decide what actions to take



Depending on the action rewards/costs, optimal policies may look differently

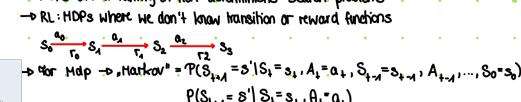
a:  $\pi(s, a)$  Deterministic policy

$a_t \sim \pi(\cdot | s_t)$  Stochastic policy



### Markov Decision Processes:

- Markovian → given present state, future and past are independent
- MDP defined by 4-tuple  $(S, A, T, R) \rightarrow S: \text{a set of states } s \in S$
- ↳ start state (distribution)
  - ↳ (optimal) terminal state
  - ↳ discount factor:  $\gamma$
- MDPs are a family of non-deterministic search problems
- RL: MDPs where we don't know transition or reward functions



→ for Mdp → Markov =  $P(S_{t+1}|S_t) = P(S_{t+1}|s_t, a_t, s_{t+1}, a_{t+1}, \dots, s_0, a_0)$

$P(S_{t+1}|s_t, a_t, s_{t+1})$



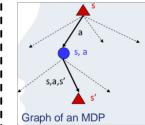
## Active Learning:

- Full reinforcement learning: don't know transitions  $T(s, a, s')$ , rewards  $R(s, a, s')$   
 can choose any action  
 Goal: learn optimal policy  
 → learner makes choices! → fundamental trade-off: exploration vs. exploitation  
 → NOT offline planning! → take actions and find out what happens

## Q-Learning:

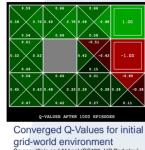
$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_a Q_k(s', a)]$

learn  $Q^*(s, a)$  values:  
 → receive sample  $(s, a, s', r)$   
 → consider old estimate:  $Q(s, a)$   
 → consider new sample estimate:  
 $\text{sample} = R(s, a, s') + \gamma \max Q(s', a')$   
 → incorporate new estimate into running average:  $Q(s, a) \leftarrow (1-\alpha)Q(s, a) + (\alpha)[\text{sample}]$



## Exploration / Exploitation:

- Several schemes for forcing exploration (simplest: random actions (ε-greedy))  
 ↳ every time step, flip a coin; probability  $\epsilon$ , act randomly; probability  $1-\epsilon$ , act according current policy
- Problems: explore space but keep thrashing around once learning done  
 ↳ solution: lower  $\epsilon$  over time or exploration functions
- Deep Q-Networks (DQN) → Q-learning with deep neural function approximator (Q-network)  
 ↳ discrete and finite set of actions  $A$ , uses ε-greedy policy to select actions



## Overview → so far:

### Things we know

if we know MDP

- compute  $V^*, Q^*, \pi^*$  exactly
- evaluate a fixed policy  $\pi$

if we don't know MDP

- we can't estimate MDP then solve
- we can estimate  $V$  for a fixed policy  $\pi$
- we can estimate  $Q^*(s, a)$  for optimal policy while executing exploration policy

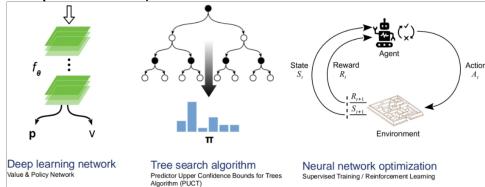
### Techniques

- ▷ Model-based DPs
  - ↳ Value and policy iteration
  - ↳ Policy evaluation
- ▷ Model-based RL
- ▷ Model-free RL
  - ↳ Value learning
  - ↳ Q-learning

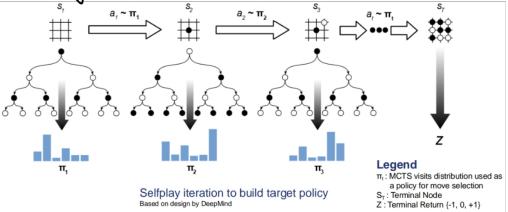
## Deep neural networks:

→ deep policy network trained to produce probability map of promising moves. Deep value network used to prune (mcmc) search tree.

### Components of AlphaZero:

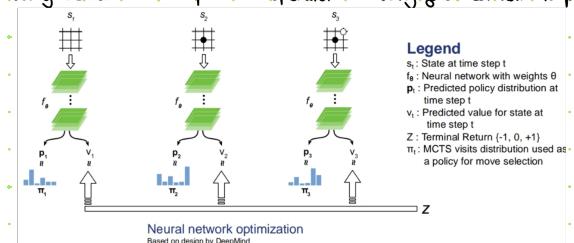


### Combining Tree Search with Neural Networks:



Generating games as training data.

### Using Search as an Improvement Operator:



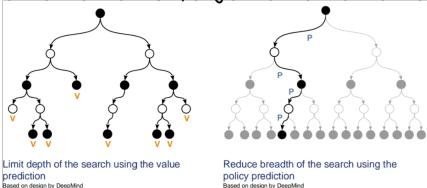
### Training Objective:

Loss formalization:  $L = \alpha(z - v)^2 - \pi^T \log p + \beta \| \theta \|^2$

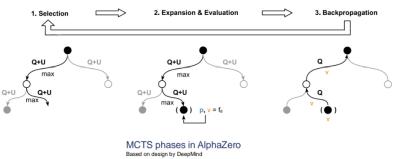
$\alpha$ : value loss factor	$\text{Mean Squared Error}$
$z$ : target value	$\text{Cross Entropy}$
$v$ : predicted value	$\text{Regularizers}$
$\pi^T$ : MCTS simulation distribution	
$p$ : policy head output	
$c$ : $L_2$ regularization constant	

## Search Principles of AlphaZero:

↳ using shared value / policy neural network



## Monte-Carlo Tree Search (MCTS):



PUCT-Algorithm: → Predictor Upper Confidence Bands for Tree Algorithm (PUCT)

→ Exploration by rollouts

→ Move selection:  $a_t = \arg\max_a(Q(s_t, a) + U(s_t, a))$ , where  $U(s_t, a) = c_{\text{puct}} P(s, a) \frac{\sqrt{N(s, b)}}{1 + N(s, a)}$

→ Update Q-values by simple moving average (SMA):  $Q'(s_t, a) \leftarrow Q(s_t, a) + \frac{1}{n} [v - Q(s_t, a)]$

$c_{\text{puct}}$ : scalar value to manage exploration

$U(s_t, a)$ : utility values

$N(s_t, a)$ : number of visits of action  $a$

$Q(s_t, a)$ : state-action values

$s_t$ : state  $s$  at time step  $t$

$a$ : action  $a$

$v$ : predicted value by neural network in  $(-1, +1)$