TECHNISCHE
UNIVERSITÄT
DARMSTADT

Fachbereich
Informatik

hessian.AI

AI101

Lecture 2: AI Systems

Recap

- **Last Week**: What is AI?

- "The question of how an intelligent being is defined, is a long and difficult one."
- Rationality behavior is about doing the right thing but is it a very good model of reality?
- We saw AIs doing great things but also how they can be tricked

Today

"How do we think about/define AI systems"

- What is an AI System?
- Environments
- Characteristics of Environments
- Agents
- The problem with Rational Agents
- Types of Agents

What is an AI System

AI System

An AI system can be defined as the study of <u>rational agents</u> and <u>their environments</u>.

Has two parts:

1. **Environment**
2. **Agent**

What is an AI System
What is an Environment

AI System
An AI system can be defined as the study of rational agents and <u>their environments</u>.

1. **Environment**
   - Definition from Oxford Dictionary:
     „environment is the surroundings or conditions in which a person, animal, or plant lives or operates".
   - In the context of AI, an environment is simply the surrounding of an (AI) agent and is where the agent operates.
   - An environment does not have to be real. It could also be artificial.

Examples
   - Selfdriving cars: Streets, traffic, weather, road signs, pedestrians, …
   - Chess: The chess boars, the chess pieces
   - ...

Characteristics of Environments (I)
Dimensions to characterize different environments

Environment	Discrete?	Observable?	Static?	Single Agent?	Accessible?	Deterministic?	Episodic?
Chess							

Characteristics of Environments (I)
Dimensions to characterize different environments

Environment	Discrete?	Observable?	Static?	Single Agent?	Accessible?	Deterministic?	Episodic?
Chess							

Discrete vs. Continuous

Has the environment (a limited/countable) number of distinct, clearly defined states then it is called <u>discrete</u>; otherwise it is <u>continuous</u>.

A self-driving car is an example of a continuous environment

Observable vs. Partially Observable or Unobservable

Is it possible to determine the complete state of the environment at each time point, then it is called <u>observable</u>; otherwise it is only <u>partially observable</u> or even unobservable.

Characteristics of Environments (II)
Dimensions to characterize different environments

Environment	Discrete?	Observable?	Static?	Single Agent?	Accessible?	Deterministic?	Episodic?
Chess	discrete	observable					

Static vs. Dynamic

If the environment does not change while an agent is acting, then it is <u>static</u>; otherwise it is <u>dynamic</u>.

Taxi driving is an example of a dynamic environment whereas Crossword puzzles are an example of a static environment.

Single Agent vs. Multiple Agents

The environment may contain <u>other agents </u>which may be of the same or different kind as that of the agent.

The agent design problems in the multi-agent environment are different from single agent environment.

Characteristics of Environments (III)
Dimensions to characterize different environments

Environment	Discrete?	Observable?	Static?	Single Agent?	Accessible?	Deterministic?	Episodic?
Chess	discrete	observable	static	Multi-agent			

Accessible vs. Inaccessible

If an agent can obtain complete and accurate information about the state's environment, then such an environment is called an <u>accessible</u> environment else it is called <u>inaccessible</u>.

An empty room whose state can be defined by its temperature is accessible / An event on earth is inaccessible.

Deterministic vs. Non-deterministic/Stochastic

If the next state of the environment is completely determined by the current state and the actions of the agent, then the environment is <u>deterministic</u>; otherwise it is <u>non-deterministic</u>.

A stochastic environment is random in nature and cannot be determined completely by an agent.

Difference between Accessible and Observable

We have an AI system, that...

- Should recognise objects in an image

- Can perceive the whole image

- Has no additional data or additional images

Is this environment observable? Is this environment accessible?

Characteristics of Environments (IV)
Dimensions to characterize different environments

Environment	Discrete?	Observable?	Static?	Single Agent?	Accessible?	Deterministic?	Episodic?
Chess	discrete	observable	static	Multi-agent	accessible	deterministic	

Episodic vs. Non-episodic/Sequential

In an episodic environment, each episode consists of the agent perceiving and then acting. The quality of its action depends just on the episode itself. In sequential environments the agent requires memory of past actions.

Why you need to know that (it is less about the correct classification)

- Different environments need different agents
- Not every algorithm works in every environment
  - e.g. some algorithm do not work with uncertainty, i.e., need perfect information

Characteristics of Environments
How to specify/define Characteristics via Questions

- Do we have a finite number of states?
- Can we observe the game perfectly?
- Are there states/actions we cannot observe?
- Is the environment changing?
- Are actions irreversible?
- Does the environment contain other agents?
- Maybe some agents have other forms than our agent?
- Do we have an opponent?
- …

- Can we observe the complete environment?
- Is our observation accurate?
- Does actions always result in the same next state?
- Do we have randomness in our environment?
- Can we determine in which state we will end taking specific actions?
- Do we need memory of the past to determine the best action?

Characteristics of Environments
Easy and Difficult Environments

Szenario 1:

Static

- We do not pay attention to possible changes in the environment

Observable

- We can at least observe our initial state

Discrete

- Possible actions can be enumerated

Deterministic

- The expected outcome of an action is always identical to the true outcome

Szenario 2:

Dynamic

- The environment is changing all the time, even without acting

Partially Observable

- Do we know our current state

Continious

- So many possiblities

Stochastically

- We can make the same action multiple times in the same state but with different results

What is more realitstic? What is easier?

Characteristics of Environments
Try it out yourself. Imagine even other environments

Environment	Discrete?	Observable?	Static?	Single Agent?	Accessible?	Deterministic?	Episodic?
Chess	Discrete	Observable	Static	Multi-Agent	Accessible	Deterministic	Sequential
Solitaire							
Poker							
Self-Driving							
Medical Diagnosis							

What is an AI Systems
What is an Agent

> **AI System**
>
> An AI system can be defined as the study of (<u>rational</u>) <u>agents</u> and their environments.

1. **Environment** ✓
2. **Agent**
   - perceives its environment (**Sense**)
   - makes decisions autonomous (**Think**)
   - acts upon the environments (**Act**)
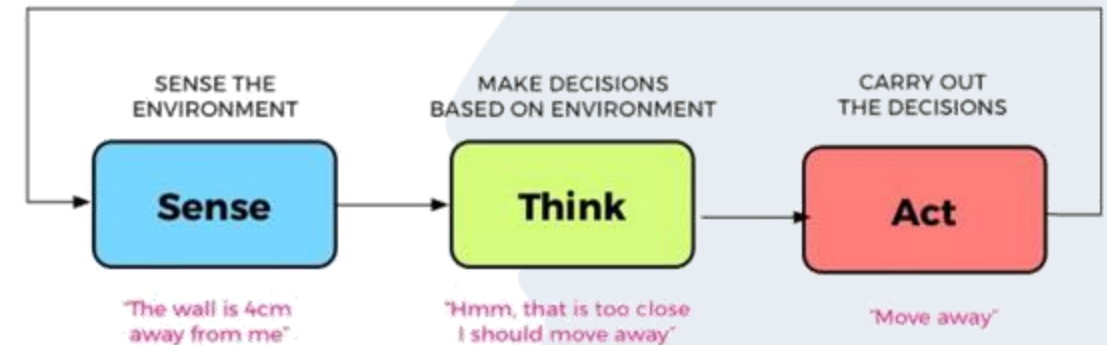
Rules of AI agents
   1. An agent must be able to perceive the environment
   2. The environments observations must be used to make decisions
   3. The decisions should result in action
   4. (The action taken by the agent must be rational)

What is an AI Systems
What is an Agent

AI System

An AI system can be defined as the study of (<u>rational</u>) <u>agents</u> and their environments.

1. Perceive

2. Decide

3. Act

Images: Pieter Abbeel and Dan Klein, CS188 Intro to AI at UC Berkley

What is an AI Systems
What is an Agent

AI System
An AI system can be defined as the study of (<u>rational</u>) <u>agents</u> and their environments.

Examples of agents

1. Humans:
    - Senses the world with their eyes, ears and other sensors
    - Thinks with their head
    - Acts with their hands, legs, vocal tract, and other actuators
2. Robots:
    - Senses the world with cameras, infrared and other sensors
    - Makes decisions based on inputs, rules or a program
    - Acts with motors for actuators
3. ...

The Problem of Rationality

An action is rational, if it maximizes the performance and yield the best positive outcome for the agent.

A rational agent is an agent that does the **right thing**

What is the "right thing"?

Be careful! Rational Agents are not omniscient

• An omniscient agent would know the actual outcome of its actions

Rational Agents are not perfect

• Rationality maximizes the expected performance

• This may not the optimal outcome

• e.g. playing in the lottery has a negative expected outcome,
  so it would be better to not play, but...

Human behavior: Irrational behavior ignores thinking, while over rational ignores feelings

The Problem of Rationality

How to measure performance

- A funtion that evaluates a sequence of actions

- Is task-dependent
 - A vacuum cleaner has a different performance criteria than a selfdriving car

General rule in designing a performance measurement

- Design the performance measure based on the desired outcome, not the desired agent behaviour

Some possible performance criteria:

Vacuum cleaner: clean floors, energy efficency, cleaning time, noise,…

Selfdriving car: Safety, comfort, time efficency, routing,…

Types of Agents
Reflex Agent

Reflex Agent

Select action on the basis of only the current percept but ignores the percept history.
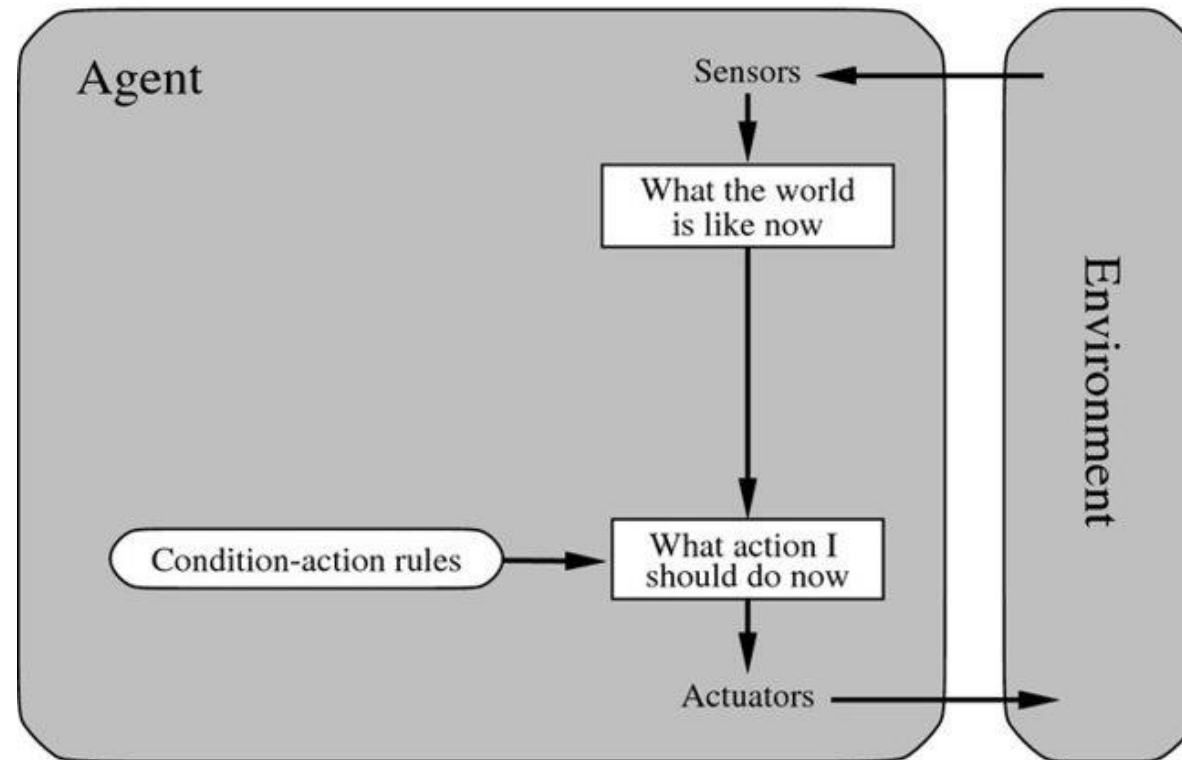
Types of Agents
Reflex Agent

Reflex Agent
Select action on the basis of only the current percept but ignores the percept history.

- Implemented through condition-action rules, i.e. "map state to action"
- Makes a very bad chess player
    - Does not look on the board, only at the last move played
    - No explicit goal of winning the game

Problem

- Very limited in ist decision making
- No knowledge about anything which the agent cannot actively perceive
- Can become very hard to handle (store, update,…) in complex environments

Types of Agents
Model-based Agent

Model-based Agent

These agents choose their actions like reflex agents do, but they have a better comprehensive view of the environment, i.e. keep track of the world state
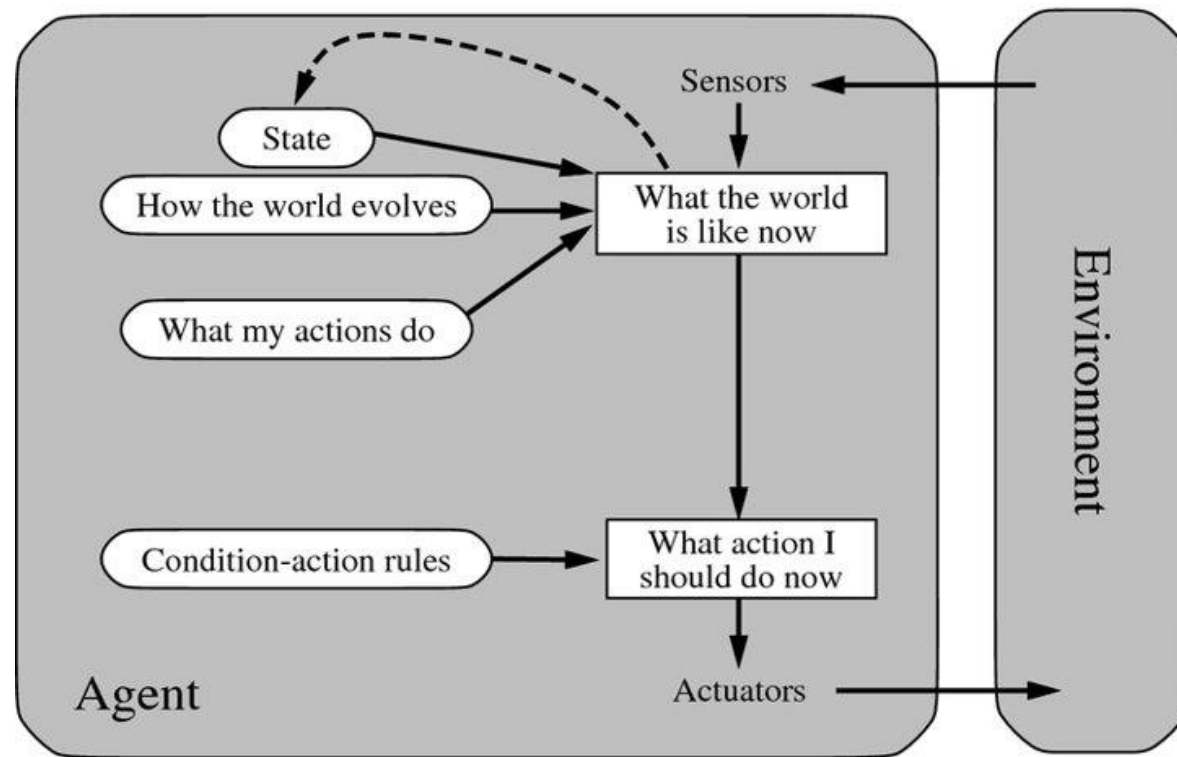
Types of Agents
Model-based Agent

Model-based Agent
These agents choose their actions like reflex agents do, but they have a better comprehensive view of the environment, i.e. keep track of the world state

- Input is not only interpreted, but mapped into an internal state (and state dynamics) description (a world model)

- Makes a better chess agent

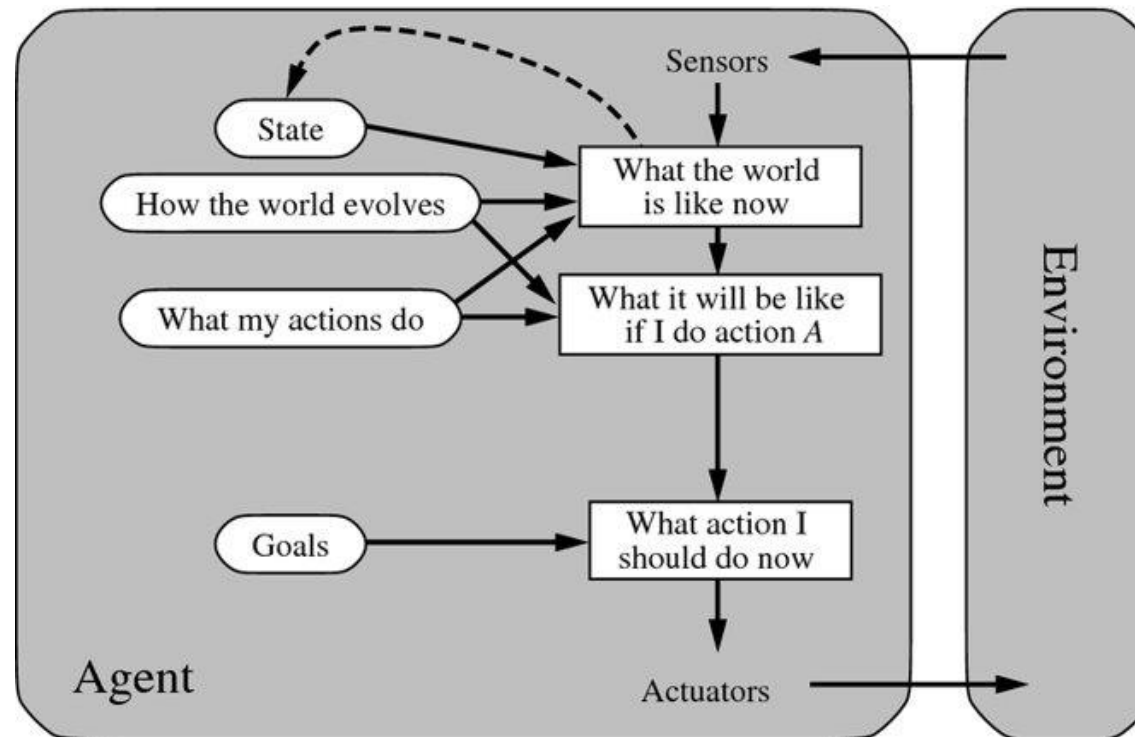- could keep track of the current board situation when its percepts are only the moves

Problem

- How does my actions affect the world?

- What world model do I desire?

Types of Agents
Goal-based Agent

Goal-based Agent

These agents build on the information that a model-based agent stores but in addition knows what states are desirable..

Types of Agents
Goal-based Agent

> **Goal-based Agent**
>
> These agents build on the information that a model-based agent stores but in addition knows what states are desirable (but not in a gradual way).

- The agent knows what states are desirable and will try to choose actions accordingly

- Main difference to previous approaches is that it takes decision-making into account
  - e.g. "What will happen if I do such-and-such?",
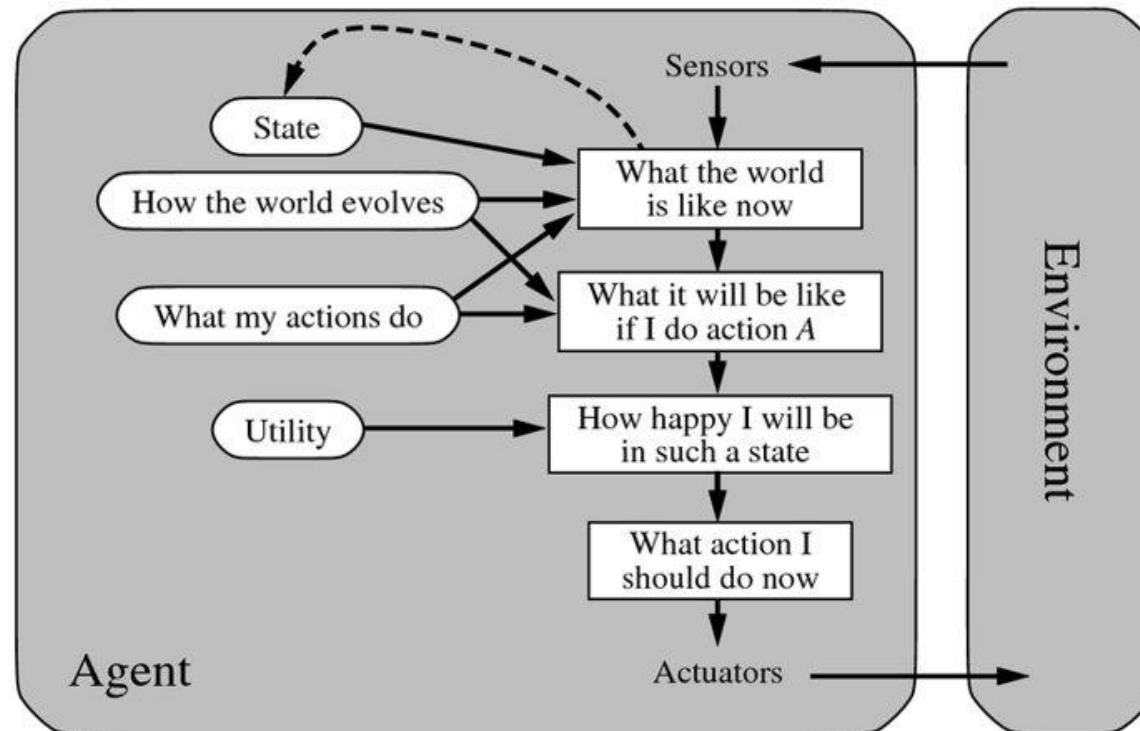 - "What will make me happy?"

Problem

- Things become difficult when long sequences of actions are required to find/achieve a goal

Types of Agents
Utility-based Agent

Utility-based Agent

Instead of providing goals, utility-based agents use a utility function for providing a way to rate each action/scenario based on the desired result

Types of Agents
Utility-based Agent

Utility-based Agent
Instead of providing goals, Utility-based agents use a utility function providing a way to rate each action/scenario based on the desired result

- Goals provide a binary distinction, while a utility function provides a contiuous scale
- Can help selecting between conflicting goals (e.g. is speed or safety more important)
- Certain goals can be reached in different ways
    - "Alle Wege führen nach Rom"
    - Some ways are quicker, safer, more reliable,… (have a higher utility)
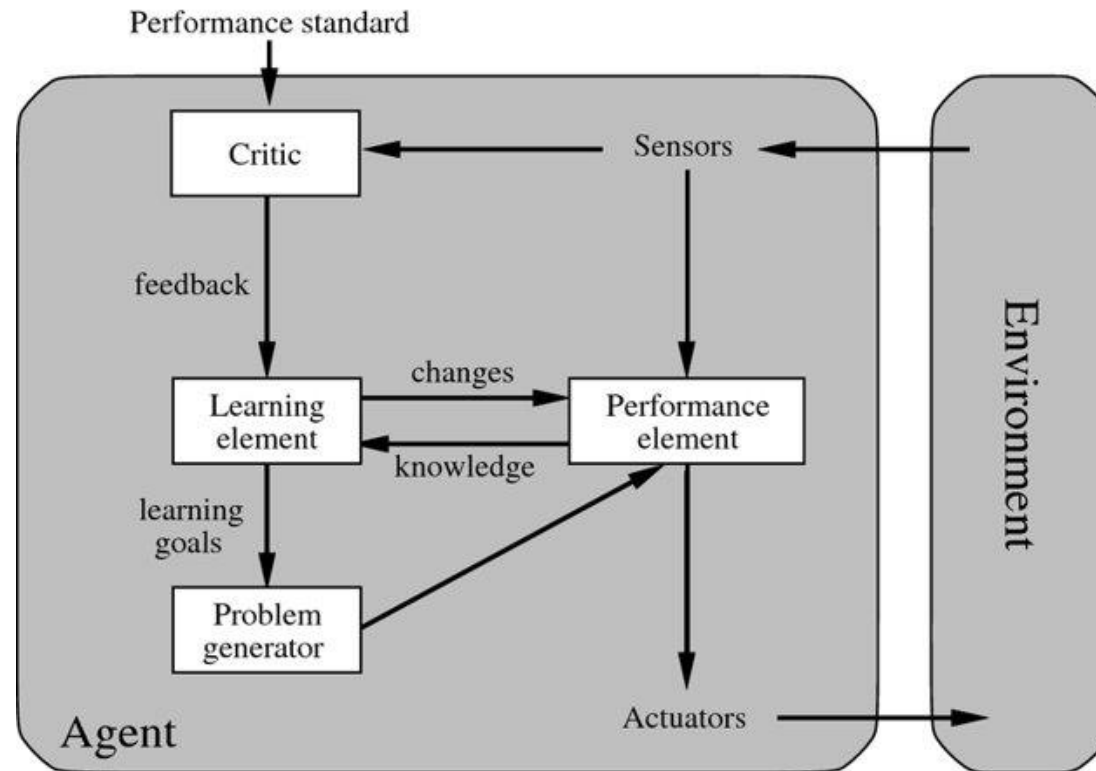
Utility function

- Maps a state or sequence of state onto a real number

Types of Agents
Learning Agent

Learning Agent

These agents employ an additional learning element to gradually improve and become more knowledgeable over time about an environment.

Types of Agents
Learning Agent

- A learning agent can learn form its past experience,
  - i.e. is able to adapt automatically based on its experience
- Is more robust toward unknown environments
- A learning agent has four conceptual components

1. **Learning Element**:
   It is responsible for making improvements by learning from the environment
2. **Critic**:
   Gives feedback, describing how well the agent is doing with respect to a fixed measurement
3. **Performance Element**:
   It is responsible for selecting actions
4. **Problem Generator**:
   Responsible for suggesting actions that will lead to new experiences

How to Make Agents Intelligent

Intelligent agents make intelligent actions
…but how do we decide what an intelligent action is
and how do we select them?

Approaches

- Search Algorithms

- Reinforcement Learning

- Genetic Algorithms

- …

How to Make Agents Intelligent
Search Algorithms

- Understand/Define "finding a good action" as a search problem and use search algorithms

- Spoiler: Most common search algorithms are tree-based
  - Bread-First is none of them

Topic of the next lectures

How to Make Agents Intelligent
Reinforcement Learning

Conditioning of Animals
Designed by slidesgo (freepik)

Pain Signal when Touching a Hot Plate
Designed by macrovector and brgfx (freepik)

- Developed within the field of psychology

- Trial and Error

- Reactions/Actions are based on our observation and experience

How to Make Agents Intelligent
## Reinforcement Learning Loop

State
S_t

Reward
R_t

Agent

Action
A_t

$\dfrac{R_{t+1}}{S_{t+1}}$

Environment

Interaction Loop between Agent & Environment

Part of the ML block (later this semester)

How to Make Agents Intelligent
Genetic Algorithms

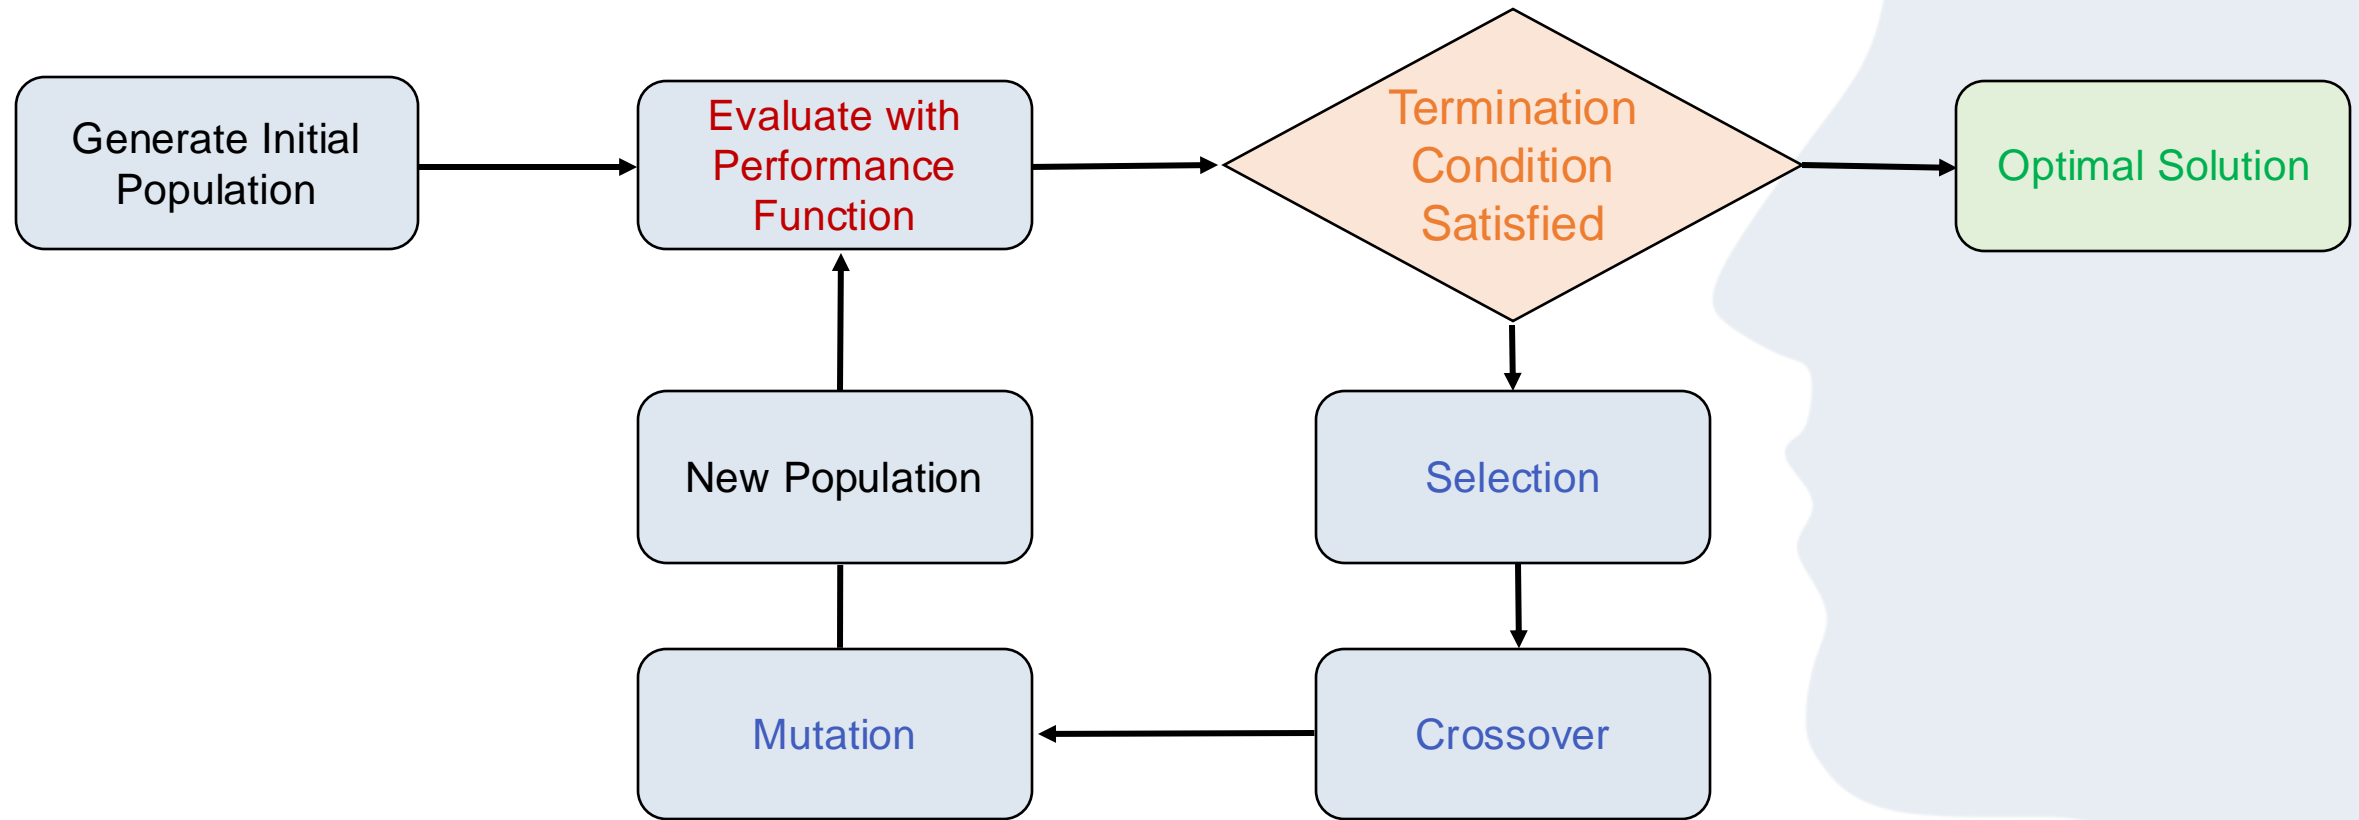Survival of the fittest inspired by Charles Darwin's theory of natural evolution

https://www.youtube.com/watch?v=Aut32pR5PQA

How to Make Agents Intelligent
Genetic Algorithms

Survival of the fittest inspired by Charles Darwin's theory of natural evolution

How to Make Agents Intelligent
Genetic Algorithms

Why is sexual reproduction so prevalent in nature when asexual reproduction seems far more efficient at promoting fitness?

A Livnat, C Papadimitriou, J Dushoff, MW Feldman, A mixability theory for the role of sex in evolution. *Proc Natl Acad Sci USA* **105**, 19803–19808 (2008)

A Livnat, C Papadimitriou, N Pippenger, MW Feldman, Sex, mixability, and modularity. *Proc Natl Acad Sci USA* **107**, 1452–1457 (2010)

Christos as featured in LogiComix. Image courtesy of Apostolos Doxiadis, Christos Papadimitriou, Alecos Papadatos, and Annie di Donna, 2009, Logicomix Print Ltd. and Bloomsbury Publishing Plc.

How to Make an Agent Intelligent
How to Improve the Performance

How do humans improve?

- We perform a task

- We remember the results

- We learn based on experience

- We go back to step 1

How do we improve agents?

- We perform a task

- We save the results of the attempt

- We adjust based on the saved data

- We go back to step 1

Is there a difference?

Training in Machine Learning (Lecture 9+10) is very similar to this idea

Summary

- Agents and environments are the main building blocks of AI software
- An agent is an entitiy that takes decisions based on what it perceives from its environment
- There are several types of agents based on their degree of perceived intelligence
- An ideal agents always tries to maximize its performance

You should be able to:

Explain the fundamental structure of an agent and the difference in agent types

Given a description, identify agents and environments as well as their properties

Next Week: Problem-solving, Uninformed Search, Informed Search